
Sharper Boundaries: Position-Aware CRF Emissions and an Embedding Adapter for Predicting Proteolytic Peptides and Propeptides

Anonymous Author(s)

Affiliation

Address

email

Abstract

Proteolysis cuts precursor proteins into mature peptides and propeptides, and predicting where is a sequence segmentation task with direct applications in proteomics, vaccine design, and disease research. The strongest existing approach couples frozen protein language model (pLM) embeddings with a CNN-BiLSTM encoder and an extended-state linear-chain conditional random field (CRF) decoder, but its CRF emissions are identical for every state of a given segment type at a given position, even though the decoder’s own state structure already distinguishes a segment’s start, interior, and end. We close this gap with two small, independent additions: a zero-initialized boundary head that adds position-specific corrections to the relevant CRF emissions, and a lightweight adapter that re-projects the pLM embedding before the encoder. Evaluated with the same 5×4 nested cross-validation DeepPeptide itself uses, on a rebuilt UniProtKB/Swiss-Prot 2026 dataset, both additions outperform the base architecture on ESM-2 embeddings individually ($+0.021$ and $+0.026$ F1) and combine to a $+0.054$ F1 gain over the 0.576 ± 0.029 base, slightly more than the sum of the two. On higher-capacity ESM-C 6B embeddings the base architecture is statistically indistinguishable from its ESM-2 counterpart under the same protocol (0.588 ± 0.016), yet the boundary head is worth two and a half times as much there and the two additions together reach 0.666 ± 0.018 : capacity the base decoder cannot use is capacity the additions can.

1 Introduction


Proteomic databases record intact sequences, not what proteases leave behind. A multi-epitope vaccine construct can therefore be checked at the design stage for the fragments it will yield in vivo, and proteolysis is itself disrupted in neurodegenerative, oncological and endocrine disease, where expected and observed cleavage patterns can be compared.

The benchmark for this task rewards detection, not localization. A predicted segment counts as correct if both of its endpoints fall within ± 3 residues of a ground-truth segment of the same type, and a single F1 score at this tolerance is reported (Figure 1). Segments are only 5–50 residues long, so on a short peptide this tolerance is comparable to the length of the segment itself. A model can therefore raise its F1 simply by finding more segments while its boundaries stay imprecise, and conversely it can place boundaries much more precisely and gain almost nothing, because predictions that close already counted as correct. The second case is our concern: the metric cannot see the improvement we are after.

We therefore make boundary localization the target of both the model and the evaluation. Our baseline is DeepPeptide [Teufel et al., 2023b], the strongest general-purpose model for this formulation.

Submitted to the AI for Drug Discovery workshop (NeurIPS 2026).

(a) baseline cuts fall inside the ± 3 window, not on the residue



(b) where they attach





Figure 1: The problem, the two blocks and what they buy. (a) One held-out precursor, with the ± 3 acceptance window shaded around each annotated cleavage site. Both of the baseline’s cuts fall inside the window, so they count as found at the headline tolerance and not at an exact match, while the model with both additions lands on the annotated residue. (b) Where the boundary head and the adapter attach to the base pipeline. (c) The F1 gap to the ESM-2 base at each tolerance, flat only if a curve were a parallel translation of the baseline. Corrected matcher, error bars over the five outer folds.

It has three parts: a frozen protein language model, a CNN-BiLSTM encoder, and a linear-chain CRF. The CRF has separate states for the start, the interior, and the end of a segment, so any valid path must visit them in that order. The decoder is therefore boundary-aware by design. The encoder, however, gives it nothing to work with: every state of the same segment type receives the same emission score. Closing that gap is the natural way to sharpen boundaries rather than merely find more segments. We close it with two small, independent additions and evaluate them under 5×4 nested cross-validation on the DeepPeptide dataset rebuilt from the 2026 UniProtKB/Swiss-Prot release (Sections 4.1 and 5).

Our contributions are the two blocks, the decomposition of what each buys under nested cross-validation, and a quantified account of why this benchmark resists a cleaner protocol: homology-aware folds are not exchangeable, their spread exceeds either addition measured on its own, and averaging it away costs roughly 200 GPU-hours per configuration.

2 Related work

Most predictors of proteolytic cleavage are specific by construction: one line covers a fixed set of proteases with known recognition motifs, one enzyme or family at a time [Duckert et al., 2004, Song et al., 2012, Li et al., 2023, 2020a,b], a second targets one product class in one kind of organism, most often neuropeptides, and scores cleavage only near basic residues [Southey et al., 2006, Wang et al., 2024]. Neither answers what an organism’s full complement of proteases does to a given precursor. General-purpose models train a compact head on frozen protein-language-model embeddings [Lin et al., 2023, ESM Team, 2024], but mostly score peptides that have already been excised [Du et al., 2024, Zhu et al., 2025] rather than locating them, and the pre-pLM PeptideLocator [Mooney et al., 2013] returns a per-residue heatmap rather than a segmentation. DeepPeptide [Teufel et al., 2023b] is, to our knowledge, the only model that segments a precursor into typed peptide and propeptide spans, and it established the homology-partitioned benchmark used here [Teufel et al., 2023a], so we take both its architecture and its data pipeline as our starting point (Appendix A).

61 3 Method

62 **The base architecture.** DeepPeptide passes a frozen ESM-2 embedding through a CNN-BiLSTM
63 stack and decodes it with a linear-chain CRF that expands each of {Peptide, Propeptide} into a
64 chain of up to 50 position states, 101 in all, so a legal path can only realize a contiguous segment of
65 admissible length. The encoder computes one emission per label and shares it across all 50 states of
66 that label: the decoder is boundary-aware, the features feeding it are not. The two additions close
67 that gap from opposite ends of the pipeline (Figure 1b). Other embedding sources and architectural
68 modifications were screened as well, with verdicts in Appendix D.

69 **Boundary head.** A small feed-forward block (LayerNorm, Linear, GELU, Linear, hidden size 64)
70 is applied to the contextual per-residue features. For each position and each segment type it emits
71 three numbers, scoring how much that position looks like a segment start, an interior position, or a
72 segment end, and these are added to the emissions of the corresponding position-fixed CRF states.
73 The final linear layer is initialized to zero, so training begins at exactly the base model and the head
74 only learns corrections that lower the loss. It costs 4,678 parameters on a trainable stack of 224,710.

75 **Adapter.** The second addition leaves the decoder alone and acts on the input. Before the per-
76 residue embedding reaches the CNN-BiLSTM it passes through LayerNorm, Linear, GELU, which
77 re-projects it and reduces its width (1280 to 256 for ESM-2). A pLM embedding is trained on
78 masked-residue recovery rather than cleavage-site localization, and its feature distribution suits nei-
79 ther this task nor the small head that consumes it, so a trainable re-projection in front of a frozen
80 encoder is the standard remedy. It is the more expensive of the two, a net +232,704 parameters.

81 4 Data and evaluation

82 4.1 Dataset

83 A precursor $x = (a_1, \dots, a_L)$ is labelled per residue with $y_t \in \{\text{None}, \text{Peptide}, \text{Propeptide}\}$
84 and contiguous runs of a label form typed segments, so the task is to recover which stretches of a
85 precursor become mature peptides and which become propeptides that are excised and discarded.
86 DeepPeptide built its dataset from PEPTIDE and PROPEP annotations in the 2022 Swiss-Prot release.
87 We rebuilt it with the same pipeline on the 2026 release [UniProt Consortium, 2025]. The collection
88 grows from 8,449 proteins to 9,619, of which 8,897 both carry ESM-2 embeddings and enter the
89 five folds used here. Folds come from GraphPart [Teufel et al., 2023a] at a 30% pairwise-identity
90 ceiling, balanced by cleavage-motif class. Segment-length filtering, motif balancing and the full
91 composition of the rebuild are given in Appendix B.

92 4.2 Evaluation criterion

93 Peptides and propeptides are matched separately. A prediction is a true positive when its start and its
94 end both fall within $\pm\tau$ residues of a true segment of the same type. Unmatched predictions are false
95 positives and unmatched true segments false negatives. Following the original evaluation we keep
96 $\tau = 3$ as the headline setting, since annotated sites carry a few residues of experimental uncertainty.
97 Sweeping τ down to zero then asks a second question, not whether a peptide was found but how
98 precisely its ends were placed. The reference implementation of this criterion carries an upstream
99 variable-shadowing bug that marks the wrong true segment as matched, so everything reported here
100 is re-scored with a corrected matcher (Appendix C).

101 **Agreement with the published baseline.** DeepPeptide reports precision 0.68 and recall 0.49 at
102 ± 3 under nested cross-validation on the 2022 data, an implied F1 of 0.570. Running the base
103 architecture on that same release under the same protocol isolates the release from every other dif-
104 ference. It scores 0.574 ± 0.069 with the original matcher, directly comparable to the published
105 figure, and 0.590 ± 0.064 with the corrected matcher of Appendix C. On the 2026 rebuild it reaches
106 0.576 ± 0.029 (Table 1), so neither the rebuild nor our reimplement moved the baseline.

Table 1: 5×4 nested cross-validation with the corrected matcher, mean $\pm$ std. over the five outer folds. Precision and recall are at the headline ± 3 tolerance, and the four F1 columns tighten the boundary-match tolerance to an exact residue. *Growth* is how much a row’s F1 advantage over the ESM-2 base widens between ± 3 and an exact match, so a value near zero would mean its tolerance curve is a parallel translation of the baseline.

Boundary	Adapter	at ± 3		F1 by tolerance				growth
		P	R	± 3	± 2	± 1	exact	
<i>ESM-2</i>								
—	—	0.621 ± 0.025	0.539 ± 0.045	0.576 ± 0.029	0.544	0.467	0.354	—
✓	—	0.673 ± 0.029	0.538 ± 0.036	0.597 ± 0.026	0.566	0.488	0.384	+0.010
—	✓	0.628 ± 0.019	0.579 ± 0.040	0.602 ± 0.026	0.572	0.491	0.396	+0.016
✓	✓	0.667 ± 0.032	0.598 ± 0.026	0.630 ± 0.021	0.599	0.518	0.432	+0.023
<i>ESM-C 6B</i>								
—	—	0.620 ± 0.017	0.560 ± 0.027	0.588 ± 0.016	0.558	0.483	0.371	+0.005
✓	—	0.731 ± 0.010	0.572 ± 0.030	0.641 ± 0.022	0.608	0.532	0.429	+0.010
—	✓	0.598 ± 0.012	0.607 ± 0.031	0.602 ± 0.017	0.573	0.497	0.394	+0.014
✓	✓	0.690 ± 0.013	0.646 ± 0.033	0.666 ± 0.018	0.636	0.554	0.466	+0.021

5 Results

Development ran in two stages: screening and confirmation.

Screening. An earlier split divided the data into seven GraphPart folds with fixed roles: four for training, one for epoch selection, one for comparing architectures and one held back. Roughly a dozen modifications were screened against it. Appendix D gives the protocol, the full verdict table and the per-candidate figures.

The folds of that split are not interchangeable (Figure 3), and under one assignment of roles several modifications changed the *sign* of their measured effect between the two held-out folds. A single draw of this kind cannot resolve an effect of 0.02–0.03, so we report none of these numbers as findings. What it can do is separate candidates that are consistently unhelpful from candidates worth paying for, and two, the boundary head and the adapter, sat at or above the base throughout.

Confirmation. The two survivors are re-evaluated with the 5×4 nested cross-validation DeepPep-tide itself uses, on ESM-2 and on ESM-C 6B, the embedding entering as a factor of the design rather than as a third candidate. For each of five outer folds, four models are trained on three of the remaining folds and validated on the fourth, then tested on the outer fold, which they never saw. An outer fold’s score is the mean over its four models and the estimate is the mean over the five outer folds. The reported spread is the standard deviation across outer folds rather than across all 20 cells, since four models sharing an outer fold share a test set. Selection happened at the screening stage, so what follows confirms two pre-specified hypotheses rather than searching a space.

Table 1 reports the same 2×2 factorial on both embeddings under that protocol, with the corrected matcher.¹

The two additions act on different errors. Each improves F1 by a similar amount on its own, and together they give 0.054 for ESM-2 and 0.078 for ESM-C, slightly more than the sum of the parts. Splitting F1 into its components explains why. The boundary head is almost purely a precision effect: precision rises by 0.052 while recall does not move at all (−0.001). The adapter is mostly a recall effect: recall rises by 0.040 against 0.007 of precision. Applied together they recover both, 0.046 precision and 0.059 recall. They add up because they act on different error types.

Tightening the tolerance. Figure 1 and the F1 columns of Table 1 sweep τ from ± 3 to an exact match. Absolute F1 falls steeply for every model: the strongest configuration goes from 0.666 at ± 3 to 0.466 at the exact residue. Placing a cleavage site exactly is still an open problem.

A model that is better overall retains a different fraction of its ± 3 score even when no boundary is placed better, so we compare absolute gaps rather than ratios. Every gap widens as the tolerance

¹Implementation, the run configuration behind every row and a synthetic smoke test: <https://anonymous.4open.science/r/cleavage-site-segmentation-5851>

139 tightens (the *growth* column of Table 1), and each stays nearly flat from ± 3 to ± 1 before opening
140 up at the last step, so what separates these models is specifically whether they hit the exact residue.

141 A widening gap is still not proof of better localization, since a model that finds more segments also
142 gains ground at every tolerance. To hold detection fixed we compared each variant against the base
143 on the true segments that *both* localize within ± 3 , matched cell by cell and segment by segment.
144 On that set the share of boundaries placed on the exact residue rises from 0.630 to 0.652 for the
145 boundary head, from 0.619 to 0.668 for the adapter and from 0.621 to 0.698 for the two together,
146 each positive on all five outer folds, on the segment counts given in Appendix E. A gate is itself a
147 choice, so Appendix F asks the same question without one, under an overlap criterion rather than a
148 tolerance. It answers what the tolerance cannot: when a hit needs only a single residue of overlap,
149 neither addition gains anything measurable (-0.013 for the head, $+0.006$ for the adapter), and both
150 gain steadily as the overlap requirement tightens. What they buy is placement rather than coverage,
151 which is why a criterion with a ± 3 window records part of it as recall.

152 **Capacity the base decoder cannot use.** The base architecture on ESM-C 6B reaches $0.588 \pm$
153 0.016 against 0.576 ± 0.029 on ESM-2. The intervals overlap, so an embedding twice as wide
154 buys no confirmed improvement on its own, while either architectural addition on the narrower one
155 does. What the additions are worth on it, though, depends on which addition. Measured against its
156 own base and paired by outer fold, the boundary head gains $+0.054 \pm 0.026$ on ESM-C 6B against
157 $+0.021 \pm 0.007$ on ESM-2, the adapter $+0.014 \pm 0.014$ against $+0.026 \pm 0.018$, and the two together
158 $+0.078 \pm 0.009$ against $+0.054 \pm 0.016$, all six positive on five folds of five. The split follows
159 what each block does. The adapter re-projects an embedding trained for masked-residue recovery,
160 and the less mismatched that embedding, the less there is to re-project, whereas the head supplies
161 position-specific evidence, and a richer embedding carries more of it. Read the other way, the same
162 embedding swap is worth $+0.012$ F1 under the base architecture and $+0.036$ under both additions.
163 The two embeddings cover slightly different protein sets, 8,897 against 8,999, and Appendix F
164 reproduces the same split on the protein-cells common to all runs. The capacity is in the wider
165 embedding either way. What changes is whether the decoder can reach it.

166 6 Conclusion

167 A zero-initialized boundary head at the decoder and a lightweight adapter at the input improve
168 segment F1 by $+0.054$ together on ESM-2 and by $+0.078$ on ESM-C 6B. They combine because
169 they act on different errors, the head suppressing spurious segments and the adapter recovering true
170 ones the base model placed outside the window, and both place boundaries more precisely, so their
171 advantage widens as the tolerance tightens (Figure 1). The evaluation carries a second message:
172 homology-aware partitioning leaves the folds taxonomically non-exchangeable, and their spread
173 exceeds either addition measured on its own.

174 **The folds stay unequal.** Averaging over folds does not make them comparable, it only stops one
175 of them deciding the result. GraphPart keeps homologs together, and taxa *are* homology clusters: the
176 five outer folds differ in size by a factor of two (1,263 to 2,572 proteins), 313 of 714 *Conus* sequences
177 fall in fold 1 against 16 in fold 2, and all 293 *Cyriopagopus* avoid fold 0 entirely (Table 3). The base
178 architecture accordingly scores 0.530, 0.572, 0.576, 0.599 and 0.604 on the five outer folds, a range
179 of 0.074, three and a half times the $+0.021$ effect the same experiment resolves once averaged over
180 them.

181 **Cost.** Nothing was selected on outer-fold scores inside confirmation, since every cell of the 2×2
182 is reported, but the screening that chose the candidates ran on the same proteins the folds are drawn
183 from, and sealing that would need a third held-out level. One cell costs about 10 GPU-hours and
184 the full $2 \times 2 \times 2$ grid roughly 1,600. A third level multiplies that again while cutting the training
185 fraction from $3/5$ to $2/5$ of an already small dataset. At this size the protocol cannot be made both
186 clean and affordable, so we state the trade-off rather than leave it implicit.

187 **Limitations.** The spread covers fold composition, not seed variation, so the true uncertainty is
188 wider. The ideas that were only ever screened, among them a structural channel, an auxiliary bond
189 loss and LoRA fine-tuning, remain unconfirmed (Appendix D).

References

- José Juan Almagro Armenteros, Marco Salvatore, Olof Emanuelsson, Ole Winther, Gunnar von Heijne, Arne Elofsson, and Henrik Nielsen. Detecting sequence signals in targeting peptides using deep learning. *Life Science Alliance*, 2(5):e201900429, 2019. doi: 10.26508/lsa.201900429.
- Zhenjiao Du, Xingjian Ding, William Hsu, Arslan Munir, Yixiang Xu, and Yonghui Li. pLM4ACE: a protein language model based predictor for antihypertensive peptide screening. *Food Chemistry*, 431:137162, 2024. doi: 10.1016/j.foodchem.2023.137162.
- Peter Duckert, Søren Brunak, and Nikolaj Blom. Prediction of proprotein convertase cleavage sites. *Protein Engineering, Design and Selection*, 17(1):107–112, 2004. doi: 10.1093/protein/gzh013.
- Ahmed Elnaggar, Hazem Essam, Wafaa Salah-Eldin, Walid Moustafa, Mohamed Elkerdawy, Charlotte Rochereau, and Burkhard Rost. Ankh: Optimized protein language model unlocks general-purpose modelling, 2023. URL <https://arxiv.org/abs/2301.06568>.
- ESM Team. Esm cambrian: Revealing the mysteries of proteins with unsupervised learning. <https://www.evolutionaryscale.ai/blog/esm-cambrian>, December 2024.
- Thomas Hayes, Roshan Rao, Halil Akin, Nicholas J. Sofroniew, Deniz Oktay, Zeming Lin, Robert Verkuil, Vincent Q. Tran, Jonathan Deaton, Marius Wiggert, Rohil Badkundri, Irhum Shafkat, Jun Gong, Alexander Derry, Raul S. Molina, Neil Thomas, Yousuf A. Khan, Chetan Mishra, Carolyn Kim, Liam J. Bartie, Matthew Nemeth, Patrick D. Hsu, Tom Sercu, Salvatore Candido, and Alexander Rives. Simulating 500 million years of evolution with a language model. *Science*, 387(6736):850–858, 2025. doi: 10.1126/science.ads0018.
- Michael Heinzinger, Konstantin Weissenow, Joaquin Gomez Sanchez, Adrian Henkel, Milot Mirdita, Martin Steinegger, and Burkhard Rost. ProstT5: bilingual language model for protein sequence and structure. *NAR Genomics and Bioinformatics*, 6(4):lqae150, 2024. doi: 10.1093/nargab/lqae150.
- Fuyi Li, Jinxiang Chen, André Leier, Tatiana Marquez-Lago, Quanzhong Liu, Yanze Wang, Jerico Revote, A. Ian Smith, Tatsuya Akutsu, Geoffrey I. Webb, Lukasz Kurgan, and Jiangning Song. DeepCleave: a deep learning predictor for caspase and matrix metalloprotease substrates and cleavage sites. *Bioinformatics*, 36(4):1057–1065, 2020a. doi: 10.1093/bioinformatics/btz721.
- Fuyi Li, André Leier, Quanzhong Liu, Yanan Wang, Dongxu Xiang, Tatsuya Akutsu, et al. Procleave: predicting protease-specific substrate cleavage sites by combining sequence and structural information. *Genomics, Proteomics & Bioinformatics*, 18(1):52–64, 2020b. doi: 10.1016/j.gpb.2019.08.002.
- Fuyi Li, Cong Wang, Xudong Guo, Tatsuya Akutsu, Geoffrey I. Webb, Lachlan J. M. Coin, Lukasz Kurgan, and Jiangning Song. ProsperousPlus: a one-stop and comprehensive platform for accurate protease-specific substrate cleavage prediction and machine-learning model construction. *Briefings in Bioinformatics*, 24(6):bbad372, 2023. doi: 10.1093/bib/bbad372.
- Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023. doi: 10.1126/science.ade2574.
- Catherine Mooney, Niall J. Haslam, Gianluca Pollastri, and Denis C. Shields. Towards the improved discovery and design of functional peptides: common features of diverse classes permit generalized prediction of bioactivity. *PLoS ONE*, 7(10):e45012, 2012. doi: 10.1371/journal.pone.0045012.
- Catherine Mooney, Niall J. Haslam, Thérèse A. Holton, Gianluca Pollastri, and Denis C. Shields. PeptideLocator: prediction of bioactive peptides in protein sequences. *Bioinformatics*, 29(9):1120–1126, 2013. doi: 10.1093/bioinformatics/btt103.

- Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo, Myle Ott, C. Lawrence Zitnick, Jerry Ma, and Rob Fergus. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences*, 118(15):e2016239118, 2021.
- Jiangning Song, Hao Tan, Andrew J. Perry, Tatsuya Akutsu, Geoffrey I. Webb, James C. Whisstock, and Robert N. Pike. PROSPER: an integrated feature-based tool for predicting protease substrate cleavage sites. *PLoS ONE*, 7(11):e50300, 2012. doi: 10.1371/journal.pone.0050300.
- Jiangning Song et al. PROSPERous: high-throughput prediction of substrate cleavage sites for 90 proteases with improved accuracy. *Bioinformatics*, 34(4):684–687, 2018. doi: 10.1093/bioinformatics/btx670.
- Bruce R. Southey, Andinet Amare, Tyler A. Zimmerman, Sandra L. Rodriguez-Zas, and Jonathan V. Sweedler. NeuroPred: a tool to predict cleavage sites in neuropeptide precursors and provide the masses of the resulting peptides. *Nucleic Acids Research*, 34(Web Server issue):W267–W272, 2006. doi: 10.1093/nar/gkl161.
- Felix Teufel, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Magnús Halldór Gíslason, Silas Irby Pihl, Konstantinos D. Tsirigos, Ole Winther, Søren Brunak, Gunnar von Heijne, and Henrik Nielsen. SignalP 6.0 predicts all five types of signal peptides using protein language models. *Nature Biotechnology*, 40(7):1023–1025, 2022. doi: 10.1038/s41587-021-01156-3.
- Felix Teufel, Magnús Halldór Gíslason, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Ole Winther, and Henrik Nielsen. GraphPart: homology partitioning for biological sequence analysis. *NAR Genomics and Bioinformatics*, 5(4):lqad088, 2023a. doi: 10.1093/nargab/lqad088.
- Felix Teufel, Jan Christian Refsgaard, Christian Toft Madsen, Carsten Stahlhut, Mads Grønberg, Ole Winther, and Dennis Madsen. Deeppeptide predicts cleaved peptides in proteins using conditional random fields. *Bioinformatics*, 39(10):btad616, 10 2023b. ISSN 1367-4811. doi: 10.1093/bioinformatics/btad616. URL <https://doi.org/10.1093/bioinformatics/btad616>.
- UniProt Consortium. UniProt: the universal protein knowledgebase in 2025. *Nucleic Acids Research*, 53(D1):D609–D617, 2025. doi: 10.1093/nar/gkae1010. URL <https://academic.oup.com/nar/article/53/D1/D609/7902999>.
- Michel van Kempen, Stephanie S. Kim, Charlotte Tumescheit, Milot Mirdita, Jeongjae Lee, Cameron L. M. Gilchrist, Johannes Söding, and Martin Steinegger. Fast and accurate protein structure search with Foldseek. *Nature Biotechnology*, 42:243–246, 2024. doi: 10.1038/s41587-023-01773-0.
- Lei Wang, Chen Huang, Mingxia Wang, Zhidong Xue, and Yan Wang. NeuroPred-PLM: an interpretable and robust model for neuropeptide prediction by protein language model. *Briefings in Bioinformatics*, 24(2):bbad077, 2023. doi: 10.1093/bib/bbad077.
- Lei Wang, Zilu Zeng, Zhidong Xue, and Yan Wang. DeepNeuropePred: a robust and universal tool to predict cleavage sites from neuropeptide precursors by protein language model. *Computational and Structural Biotechnology Journal*, 23:309–315, 2024. doi: 10.1016/j.csbj.2023.12.004.
- M. Zhang, J. Zhou, X. Wang, et al. DeepBP: ensemble deep learning strategy for bioactive peptide prediction. *BMC Bioinformatics*, 25:352, 2024. doi: 10.1186/s12859-024-05974-5.
- Lei Zhu, Hao Sun, and Sen Yang. BPFun: a deep learning framework for bioactive peptide function prediction using multi-label strategy by transformer-driven and sequence rich intrinsic information. *BMC Bioinformatics*, 26(1):187, 2025. doi: 10.1186/s12859-025-06190-5.

Impact Statement

This paper improves a computational tool for predicting proteolytic cleavage products from protein sequence. Potential applications include interpreting proteomics data, informing vaccine and peptide-drug design, and studying disease-associated proteolysis. We are not aware of societal risks specific to this work beyond those generally associated with more accurate protein sequence analysis tools.

288 Appendix

289 A Extended review of prior work

290 Methods for predicting proteolytic peptides fall into three broad classes.

291 **Protease-specific models.** These predict cleavage only for a limited set of enzymes, most of which
292 have a known recognition motif. ProP [Duckert et al., 2004] targets cleavage by the PACE/PC family
293 in animals and plants. PROSPER [Song et al., 2012] and its successor PROSPERous [Song et al.,
294 2018] predict cleavage for one chosen enzyme at a time, as does the later ProsperousPlus [Li et al.,
295 2023]. DeepCleave [Li et al., 2020a] is restricted to caspases and matrix metalloproteinases, and
296 Procleave [Li et al., 2020b] likewise handles one enzyme at a time and additionally requires 3D
297 structural input. The shared limitation of this class is scope: in practice, one is usually interested in
298 the combined effect of all proteases active in an organism, not any single enzyme.

299 **Organism- or peptide-type-specific models.** A separate class targets neuropeptides or other bioac-
300 tive fragments specific to a given organism. NeuroPred [Southey et al., 2006] and the species-
301 agnostic DeepNeuropePred [Wang et al., 2024] only detect cleavage near a handful of basic residues,
302 while NeuroPred-PLM [Wang et al., 2023] does not localize cleavage sites at all but classifies an
303 already-excised sequence. A related group of models finds signal subsequences rather than any
304 bioactive peptide: SignalP [Teufel et al., 2022] and TargetP [Almagro Armenteros et al., 2019] iden-
305 tify the type of signal peptide, which indicates where a protein is trafficked.

306 **General-purpose models on pLM embeddings. Protein language models (pLMs)** are transform-
307 ers trained without supervision on large sequence corpora, typically by masked-residue recovery,
308 and they produce transferable per-residue embeddings: ESM-2 [Lin et al., 2023], the earlier ESM
309 line [Rives et al., 2021], ESM Cambrian [ESM Team, 2024, Hayes et al., 2025], and Ankh [El-
310 naggar et al., 2023]. Compact task-specific heads are then trained on top of these frozen embed-
311 dings. Besides **DeepPeptide** [Teufel et al., 2023b] itself, this recipe underlies DeepNeuropePred,
312 NeuroPred-PLM, and SignalP 6, as well as models that score already-excised peptides rather than
313 finding them: pLM4ACE [Du et al., 2024] predicts ACE-inhibitory activity, and BPFun [Zhu et al.,
314 2025] and DeepBP [Zhang et al., 2024] predict broader bioactivity, building on earlier pre-pLM work
315 such as PeptideRanker [Mooney et al., 2012]. Before pLMs, PeptideLocator [Mooney et al., 2013]
316 addressed a problem close to ours, but its output is a per-residue heatmap of similarity to known
317 bioactive peptides rather than an explicit segmentation, and it was the main point of comparison in
318 the original DeepPeptide paper.

319 **Baseline architecture: DeepPeptide.** Since we build directly on it throughout this paper (Section 3),
320 DeepPeptide’s architecture is worth describing in detail rather than treating as a black box. It first
321 encodes a precursor sequence with a frozen pLM (ESM-2), producing a per-residue embedding that
322 is not fine-tuned during training. This embedding is refined by a convolution, a single-layer bidirec-
323 tional LSTM and a second convolution. The published description does not fix the widths: they are
324 searched with Optuna inside the inner cross-validation loop. The configuration we train throughout
325 uses 32 convolutional filters of width 3 and an LSTM hidden size of 64 per direction, giving 224,710
326 trainable parameters for the base model. These features are consumed by a linear-chain CRF decoder
327 whose state set is extended well beyond the three raw labels {None, Peptide, Propeptide}: each
328 of the two positive labels is expanded into a chain of states long enough to represent segments up
329 to a fixed maximum length (101 states in total for the 50-residue cap used here, Section 4.1), so
330 that a legal path through the CRF can only realize a contiguous segment of valid length, with ded-
331 icated states marking its first and last positions. This makes segmentation, rather than independent
332 per-residue classification, the object the model is directly optimized for – and it is also where the
333 gap described in Section 1 lives: every state in this extended set that shares a type receives the same
334 emission from the encoder, regardless of whether it marks the start, interior, or end of a segment.

335 **Baseline dataset for this task.** Beyond its architecture, DeepPeptide also established the data re-
336 source this line of work relies on: precursor sequences from UniProtKB/Swiss-Prot annotated with
337 PEPTIDE and PROPEP feature types, partitioned into homology-aware folds with GraphPart [Teufel
338 et al., 2023a] at a 30% pairwise-identity threshold and additionally balanced by cleavage-flanking
339 motif, then evaluated with the same family of nested cross-validation we adopt (Section 5). This
340 gave the field a standardized, already homology-controlled benchmark rather than an ad hoc collec-
341 tion of previously published peptide lists, and both the curated data and the construction pipeline

that produced it are public. Section 4.1 describes what we do with that resource: rebuilding it on a current UniProtKB/Swiss-Prot release, which adds 1,170 proteins and refreshes the annotations behind every segment.

Among these, DeepPeptide is, to our knowledge, the only model that poses the task as full segmentation of the precursor into typed peptide and propeptide segments, and it is the strongest available baseline for this formulation. We therefore use its architecture and its dataset-construction methodology as the starting point for this work, rather than treating it as a system to audit: our contribution is an architectural addition to this general recipe (Section 3), evaluated under the same 5×4 nested cross-validation it uses (Section 4) on data rebuilt from a current release (Section 4.1).

B Dataset details

Table 2: Dataset composition: the 2022 release used by the original DeepPeptide against the 2026 release rebuilt here. Counts are before homology partitioning.

	2022	2026
Proteins	8449	9619
of which absent in 2022	—	1178 ($\approx 12\%$)
Peptide segments	6372	7431
Propeptide segments	8211	9140
Median peptide length (residues)	21	20




Figure 2: Composition of the rebuilt 2026 dataset: segment length, type, and genus distributions.

Two filters from the original pipeline are kept unchanged. Segments shorter than 5 or longer than 50 residues are dropped: the CRF’s state count fixes the upper bound, and a segment shorter than five residues is barely scoreable at a ± 3 tolerance. Folds are then balanced by cleavage-motif class, obtained by k -means ($k = 50$) on ESM-2 embeddings of the four residues flanking each annotated boundary, on top of the 30% pairwise-identity ceiling GraphPart enforces between folds.

C Metric implementation bug

While auditing the DeepPeptide implementation of the matching criterion (Section 4.2), we found a bug in the function that implements it (`get_counts_for_protein`), reproduced below in simplified form:

Table 3: Proteins per outer fold, and how three frequent genera distribute across them. *Cyriopagopus* and *Lycosa* are spider venoms, *Conus* a cone snail, and each is a homology cluster that GraphPart is obliged to keep intact. *Homo*, by contrast, is spread evenly (71/81/66/130/84).

	total	f0	f1	f2	f3	f4
All proteins	8897	1558	2572	1263	2025	1479
<i>Conus</i>	714	160	313	16	149	76
<i>Cyriopagopus</i>	293	0	62	45	9	177
<i>Lycosa</i>	163	5	42	10	106	0

```

361
362 for idx, row in true_df.iterrows(): # idx: true-segment index
363     true_start, true_stop = row['start'], row['stop']
364     for idx, row in pred_df.iterrows(): # idx overwritten with pred index
365         pred_start, pred_stop = row['start'], row['stop']
366         ...
367         if start_match and stop_match:
368             true_df.loc[idx, 'matched'] = True # BUG: idx is now the
369             pred_df.loc[idx, 'matched'] = True # predicted index, not the true one
370         break

```

The outer loop iterates over true segments and the inner loop over predicted segments, but both reuse the loop variable `idx`. By the time a match is found, `idx` refers to the predicted segment, so `true_df.loc[idx, 'matched']` marks the wrong true segment. The effect is not large on average (it is partly self-correcting), but it costs a true positive whenever it fires, so both recall and precision come out low, and F1 with them. Across the nested-CV grid the correction raises recall by 0.021 to 0.032 and precision by 0.005 to 0.010. The loss concentrates in proteins where the model predicts more segments than the protein actually has. This bug is present in the upstream DeepPeptide code, not something we introduced.

We did not silently patch this function in place, since that would shift absolute numbers without documenting it. Instead we implemented a separate, corrected matcher and used it to re-score the results in this paper: the earlier single-split runs reported in Appendix D, except two whose model classes no longer exist in the code and whose checkpoints cannot be rebuilt, and the full 5×4 nested-cross-validation grid in Table 1 (all 160 cells of the eight configurations, and the 20 further cells of the 2022-release reproduction), by re-running test-partition inference from each cell’s saved checkpoint on the machine that held them. Every table in the main text uses the corrected matcher, and the original numbers are kept alongside in the released run artifacts.

Re-scoring the grid also recomputed the original matcher on the same decoded segments, so each cell could be checked against the number already recorded for it. Over the 80 cells for which that check could be run, the two agree to 1.2×10^{-3} F1 at worst and to 3×10^{-7} at the median, so what remains is inference nondeterminism rather than disagreement between two deterministic functions of the same segments.

Across the seven runs whose every cell carries both matchers the correction shifts mean F1 by +0.015 to +0.022, always upward, since the bug was discarding true positives. On ESM-2 it leaves the ordering of the four configurations intact, base < boundary head < adapter < combined, before and after correction. That ordering is not robust and we build no argument on it: the boundary-head/adaptor gap is 0.005 against a fold-level standard deviation of 0.026, and per outer fold the full ordering holds in four folds of five after correction and three of five before it. On ESM-C 6B the ordering differs, base < adapter < boundary head < combined, the two middle configurations swapping places.

D Screening results details

D.1 Protocol

Data was split into seven GraphPart folds (30% identity threshold, motif-balanced as in Section 4.1), with folds assigned fixed roles: four folds for training, one for validation (epoch selection), one for model selection (comparing architectures), and one originally intended as a fully sealed test fold. This separated architecture selection from evaluation, but the evaluation itself was still a single

draw: one random assignment of roles, one pair of held-out folds. Several modifications changed the sign of their measured effect between the two held-out folds. Replacing ESM-2 with ESM-C 6B measured +0.074 on one and −0.011 on the other, ESM-C 6B against ESM-C 600M gave +0.028 and −0.021, and the net effect of the 3Di channel gave +0.015 and −0.018. Effects that kept their sign still moved by a lot: the isolated gated adapter measured −0.001 on one fold and +0.045 on the other. Figure 3 shows one reason, namely that the folds differ substantially in segment-length composition. The two held-out folds turn out to be complementary in segment-length composition, fold 2 concentrated at 10–24 residues and fold 5 bimodal, so results below pool them (model-select $\cup$ sealed, $\approx 2,300$ proteins) to reduce (but not eliminate) this instability. This pooling is why the “sealed” fold is not, in practice, a held-out test set independent of model selection.




Figure 3: Why the folds of the earlier seven-fold split are not interchangeable: (a) the segment-length profile of each fold against the profile of the whole dataset, (b) the L_1 distance between the two. Fold 4 tracks the overall distribution closely (0.07) while five of the seven folds sit above 0.23. Balancing was done on cleavage motif and homology, not on segment length, so this axis was left free. The same holds for the five-fold split used in the main text, where the imbalance is taxonomic (Table 3).

417 D.2 Summary of tested modifications

418 Table 4 summarizes the modifications tested under screening protocol, with the corrected matcher
 419 (Appendix C) applied throughout. Two of these rows were carried forward as candidates, the bound-
 420 ary head and the adapter (Section 3), and both were confirmed. The embedding swap was not a
 421 candidate but became a factor of the confirmation design, and it is the row this table reads least well:
 422 nested cross-validation puts ESM-C 6B at 0.588 ± 0.016 against ESM-2’s 0.576 ± 0.029 (Table 1),




Figure 4: Precision/recall trade-offs by segment type when adding the 3Di channel or the bond-prediction loss, single-split protocol, pooled held-out folds.

423 overlapping intervals, where this table records a confident +0.03. Everything else below remains at
 424 the confidence of a single pooled split.

Table 4: Modifications tested under the screening single-split protocol, with their measured effect on F1 (pooled held-out folds) and an informal verdict. None of these effect sizes should be compared directly to the nested-CV numbers in Table 1.

Modification	Type	Measured effect on F1	Verdict
ESM-C 6B instead of ESM-2	embedding swap	+0.03, unstable across folds	helps
Boundary head on ESM-C 6B	decoder addition	+0.05, +0.07, consistent	helps
Boundary head on ESM-2	decoder addition	≈ 0 (CI includes zero)	no effect
Structural channel (3Di)	extra input	propep. +0.02 / pep. -0.03, net trade-off	no effect
ESM-C 6B compression 2560→256	optimization	≈ 0 , 10× narrower input and 16× fewer trainable parameters	no effect
Gated adapter on ESM-C 6B (pLM re-projection)	input adapter	+0.022, CI [+0.008, +0.038]	helps
Bond-prediction auxiliary loss (on ESM-C 6B)	extra loss	pep. -0.05, net -0.015, neutral on ESM-2	harmful
Telescopic segment CRF	decoder addition	≈ 0 on ESM-2, +0.022 on ESM-C 6B, no CI available	unresolved

425 Note the apparent tension with Table 1: under this earlier protocol, the boundary head on ESM-2
 426 looked like it had no effect, while nested cross-validation resolved a confirmed +0.021 F1 gain
 427 for the same modification on the same embedding. We read this as evidence that the single-split
 428 protocol lacked the statistical power to resolve an effect of this size, not as a contradiction – and as
 429 the clearest illustration of why we do not treat single-split numbers as findings in this paper. Whether
 430 the boundary head’s effect is specifically larger on richer embeddings (as the ESM-C 6B rows above
 431 suggest) is accordingly left as an open question, not stated as a conclusion here.

432 D.3 Additional exploratory modifications

433 A larger set of ideas was tried under the same single-split protocol and did not show a clear, consis-
 434 tent benefit. We list them for completeness and as candidates for future re-testing, without effect-size
 435 claims:

- **Known-peptide dictionary.** An Aho–Corasick index of peptides from the training data was used to flag substring matches in a query sequence, injected as a bonus on the CRF state emissions, as a bias on the start, inside and end emissions specifically, fused into the encoder hidden state, and concatenated with the embedding before the encoder. No consistent benefit was observed, and the checkpoint for one of the variants can no longer be loaded.
- **Structural features.** Features extracted from predicted 3D structures (AlphaFold2 representations via the AFToolkit codebase, with several confidence-based filtering variants) and from the ProstT5 [Heinzinger et al., 2024] structural alphabet (3Di, as used in Foldseek [van Kempen et al., 2024]) were tested as additional input channels, and results were mixed (Table 4).
- **LoRA fine-tuning.** Low-rank adaptation of the last few pLM layers, instead of fully frozen embeddings, scored 0.558 and 0.567 against 0.621 for the frozen baseline, all three on the 2022 release under its own five-fold split rather than the protocol of Appendix D. It also ran for fewer epochs in the same time budget, though we did not measure the cost directly.
- **Projector variants.** Multi-scale and multi-branch projectors between the embedding and the CNN–BiLSTM were tried as alternative re-projection schemes to the adapter of Section 3.
- **Structural-projection width and telescopic CRF.** The width of the structural-feature projection (16/32/48 units) left pooled F1 flat at 0.691, 0.697 and 0.693, the telescopic segment CRF, an alternative decoder formulation that scores whole segments through a relative-position head, came out level with the base decoder on ESM-2 and +0.022 ahead on ESM-C 6B. Per-protein outputs were not retained for either run, so neither number carries a confidence interval and we treat the modification as untested rather than as having no effect.

None of these modifications showed a consistent gain large enough, under the single-split protocol, to justify testing under nested cross-validation ahead of the two reported in the main text. The complete experiment log (including runs not summarized above) is maintained in the project repository rather than reproduced here, since it was collected under a protocol we no longer treat as sufficient evidence on its own.

E Complete tolerance sweep

Table 5 repeats the tolerance sweep of Table 1 with the fold-level standard deviation on every cell, which the main table carries only on the headline ± 3 column for want of width. The paired localization comparison of Section 5 runs on the true segments that a variant and the base both place within ± 3 : 27,013 segments for the boundary head, 29,277 for the adapter and 29,394 for the two together, pooled over the twenty cells of each configuration.

Table 5: Segment F1 by boundary-match tolerance, 5×4 nested cross-validation, corrected matcher, mean $\pm$ std. over the five outer folds.

Configuration	± 3	± 2	± 1	exact	gap growth
ESM-2, base	0.576 ± 0.029	0.544 ± 0.027	0.467 ± 0.020	0.354 ± 0.010	+0.000
ESM-2 + boundary head	0.597 ± 0.026	0.566 ± 0.024	0.488 ± 0.015	0.384 ± 0.015	+0.010
ESM-2 + adapter	0.602 ± 0.026	0.572 ± 0.024	0.491 ± 0.011	0.396 ± 0.020	+0.016
ESM-2 + boundary head + adapter	0.630 ± 0.021	0.599 ± 0.023	0.518 ± 0.016	0.432 ± 0.013	+0.023
ESM-C 6B, base	0.588 ± 0.016	0.558 ± 0.013	0.483 ± 0.007	0.371 ± 0.013	+0.005
ESM-C 6B + boundary head	0.641 ± 0.022	0.608 ± 0.026	0.532 ± 0.035	0.429 ± 0.030	+0.010
ESM-C 6B + adapter	0.602 ± 0.017	0.573 ± 0.019	0.497 ± 0.016	0.394 ± 0.031	+0.014
ESM-C 6B + boundary head + adapter	0.666 ± 0.018	0.636 ± 0.015	0.554 ± 0.016	0.466 ± 0.021	+0.021

F Segment quality without a tolerance

The ± 3 criterion of Section 4.2 is the benchmark’s, and Section 1 says why it cannot see the improvement this paper is about. This appendix scores the same predictions without it. Predicted and true segments of the same type are matched greedily one to one by intersection over union, and a




Figure 5: F1 (± 3) against the number of training proteins, single-split protocol. ESM-2 rises from 0.498 at 40% of the training folds to 0.583 at 85%, then falls back to 0.572 at 100%, so it is still gaining over most of the range but not at the last point. The ESM-C 6B curves are flat from the smallest size tested, with movement smaller than the bootstrap interval, so nothing is observed saturating: they simply never climb.

match counts as a hit when the overlap clears a threshold τ . Precision and recall use all predicted and all true segments as their denominators, so nothing is conditioned on what either model happened to find. Every configuration is scored on the 35,504 protein-cells common to all runs, which takes the coverage difference between the two embeddings (Section 4.1) out of the comparison. On that common set the ESM-2 to ESM-C 6B swap is worth +0.005 F1 at $\text{IoU} \geq 0.5$ under the base architecture, on two outer folds of five, and +0.024 under both additions, on five of five, reproducing the split that Section 5 measures at ± 3 .

Table 6: Segment F1 under overlap matching. The first column requires only that a predicted and a true segment of the same type share one residue, and the rest require an intersection over union of at least τ . Same cells as Table 1, restricted to the protein set common to every run, mean over the five outer folds.

	overlap	$\tau \geq 0.5$	$\tau \geq 0.8$	$\tau \geq 0.9$	$\tau \geq 0.95$
<i>ESM-2</i>					
base	0.729	0.667	0.569	0.487	0.411
+ boundary head	0.715	0.668	0.585	0.515	0.438
+ adapter	0.734	0.686	0.590	0.516	0.441
+ both	0.746	0.703	0.616	0.546	0.474
<i>ESM-C 6B</i>					
base	0.738	0.672	0.574	0.496	0.419
+ boundary head	0.735	0.704	0.624	0.555	0.478
+ adapter	0.729	0.681	0.587	0.514	0.440
+ both	0.759	0.727	0.649	0.582	0.507

Placement, not coverage. The looser the criterion, the less either addition is worth. Paired by outer fold against the ESM-2 base, at a gate that asks only for a single residue of overlap the boundary head is worth -0.013 ± 0.026 and the adapter $+0.006 \pm 0.012$, neither distinguishable from zero. At $\tau \geq 0.5$ they are worth $+0.001 \pm 0.021$ and $+0.019 \pm 0.017$, and at $\tau \geq 0.9$ $+0.028 \pm 0.004$ and $+0.028 \pm 0.015$, both on five folds of five. Underneath that near-zero F1 at the overlap gate the two move in opposite directions: the head issues 219 fewer segments per fold, buying +0.022 precision for -0.035 recall, while the adapter issues 130 more and buys +0.024 recall for -0.021 precision.

Neither finds appreciably more of the annotation than the base architecture does. What changes is where the boundaries land. This is the same conclusion Table 1 reaches through a tolerance, reached without one.

The mean IoU taken over *all* true segments moves the other way for the head, 0.576 to 0.559, while its F1 rises at every threshold. That is the precision effect of Table 1 seen from another angle: the head withdraws weak predictions, so some true segments lose a poor match altogether, and the ones that keep a match are placed better.

Which end moves. Scoring cleavage sites rather than segments, at exact placement and against all annotated sites, separates the two additions again. The boundary head improves the C-side of a segment, $+0.028 \pm 0.006$ on five folds of five, and leaves the N-side alone, -0.007 ± 0.020 on two of five. The adapter improves both, $+0.022 \pm 0.006$ and $+0.028 \pm 0.009$, and the two together improve both further, $+0.042 \pm 0.010$ and $+0.060 \pm 0.006$. Split by segment type, the head's C-side gain is largest on propeptides ($+0.037 \pm 0.021$, five of five), whose C-terminal cut is where the mature peptide begins and is the weaker of the two C-sides: the base model scores 0.476 there against 0.525 on peptide C-sides.

F.1 Supplementary figures

The figures below are from the same screening stage as Table 4 and none of them should be read as confirmed evidence in the sense of Section 5.




Figure 6: Screening comparison on the pooled held-out folds at ± 3 : F1, precision and recall with bootstrap confidence intervals over 2,322 proteins. The top row is a gated-projector control at full width, not the configuration carried into the main text.




Figure 7: Boundary head × embedding interaction under the screening protocol: F1 gain from adding the boundary head on top of ESM-2 against on top of ESM-C 6B. This is the earlier, unconfirmed version of the interaction that Section 5 measures under nested cross-validation.




Figure 8: Recall on held-out peptides as a function of (a) how well-represented the protein’s genus is in training and (b) maximum sequence identity to a training segment, screening protocol. For the two ESM-C models recall tracks maximum identity to a training segment far more than genus abundance, while for the ESM-2 baseline the two axes are closer to comparable. The x axis in (a) counts training *segments* in the genus, not proteins.