
Sharper Boundaries: Position-Aware CRF Emissions and an Embedding Adapter for Predicting Proteolytic Peptides and Propeptides

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Proteolysis cuts precursor proteins into mature peptides and propeptides, and pre-
2 dicting where is a sequence segmentation task with direct applications in pro-
3 teomics, vaccine design, and disease research. The strongest existing approach
4 couples frozen protein language model (pLM) embeddings with a CNN-BiLSTM
5 encoder and an extended-state linear-chain conditional random field (CRF) de-
6 coder, but its CRF emissions are identical for every state of a given segment type
7 at a given position, even though the decoder’s own state structure already distin-
8 guishes a segment’s start, interior, and end. We close this gap with two small, inde-
9 pendent additions: a zero-initialized boundary head that adds position-specific cor-
10 rections to the relevant CRF emissions, and a lightweight adapter that re-projects
11 the pLM embedding before the encoder. Evaluated with the same 5×4 nested
12 cross-validation DeepPeptide itself uses, on a rebuilt UniProtKB/Swiss-Prot 2026
13 dataset, both additions outperform the base architecture on ESM-2 embeddings
14 individually ($+0.021$ and $+0.026$ F1) and combine to a $+0.054$ F1 gain over the
15 0.576 ± 0.029 base, slightly more than the sum of the two. On higher-capacity
16 ESM-C 6B embeddings the base architecture is statistically indistinguishable from
17 its ESM-2 counterpart under the same protocol (0.588 ± 0.016), yet the boundary
18 head is worth two and a half times as much there and the two additions together
19 reach 0.666 ± 0.018 : capacity the base decoder cannot use is capacity the additions
20 can.

21 1 Introduction

22 Proteomic databases record intact sequences, not what proteases leave behind, so the products of
23 cleavage have to be predicted. A multi-epitope vaccine construct can then be checked at the design
24 stage for the fragments it will yield in vivo, and since proteolysis is itself disrupted in neurodegener-
25 ative, oncological and endocrine disease, expected and observed cleavage patterns can be compared
26 between healthy and diseased tissue.

27 The benchmark for this task rewards detection, not localization. A predicted segment counts as cor-
28 rect if both of its endpoints fall within ± 3 residues of a ground-truth segment of the same type, and a
29 single F1 score at this tolerance is reported (Figure 1). Segments are only 5–50 residues long, so on a
30 short peptide this tolerance is comparable to the length of the segment itself. A model can therefore
31 raise its F1 simply by finding more segments while its boundaries stay imprecise, and conversely it
32 can place boundaries much more precisely and gain almost nothing, because predictions that close
33 already counted as correct. The second case is our concern: the metric cannot see the improvement
34 we are after.




Figure 1: The problem, the two blocks and what they buy. (a) One held-out precursor, with the ± 3 acceptance window shaded around each annotated cleavage site. Both of the baseline’s cuts fall inside the window, so they count as found at the headline tolerance and not at an exact match, while ours, carrying both additions, lands on the annotated residue. (b) Where the boundary head and the adapter attach to the base pipeline. (c) The F1 gap to the ESM-2 base at each tolerance, flat only if a curve were a parallel translation of the baseline. Corrected matcher, error bars over the five outer folds.

We therefore make boundary localization the target of both the model and the evaluation. Our baseline is DeepPeptide [Teufel et al., 2023b], the strongest general-purpose model for this formulation. It has three parts: a frozen protein language model, a CNN-BiLSTM encoder, and a linear-chain CRF. The CRF has separate states for the start, the interior, and the end of a segment, so any valid path must visit them in that order. The decoder is therefore boundary-aware by design. The encoder, however, gives it nothing to work with: every state of the same segment type receives the same emission score. Closing that gap is the natural way to sharpen boundaries rather than merely find more segments. We close it with two small, independent additions and evaluate them under 5×4 nested cross-validation on the DeepPeptide dataset rebuilt from the 2026 UniProtKB/Swiss-Prot release (Sections 4.1 and 5).

Our contributions are the two blocks, the decomposition of what each buys under nested cross-validation, and a quantified account of why this benchmark resists a cleaner protocol: homology-aware folds are not exchangeable, their spread exceeds either addition measured on its own, and averaging it away costs roughly 200 GPU-hours per configuration.

2 Related work

Most predictors of proteolytic cleavage are specific by construction. One line covers a fixed set of proteases with known recognition motifs, one enzyme or family at a time [Duckert et al., 2004, Song et al., 2012, Li et al., 2023, 2020a,b], a second scores cleavage near basic residues for one product class, usually neuropeptides [Southey et al., 2006, Wang et al., 2024], and neither answers what an organism’s full complement of proteases does to a precursor. General-purpose models train a compact head on frozen pLM embeddings [Lin et al., 2023, ESM Team, 2024] but mostly score peptides already excised [Du et al., 2024, Zhu et al., 2025] rather than locating them, and PeptideLocator [Mooney et al., 2013] returns a per-residue heatmap rather than a segmentation. DeepPeptide [Teufel et al., 2023b] is, to our knowledge, the only model that segments a precursor into typed peptide and propeptide spans, and it established the homology-partitioned benchmark used here [Teufel et al., 2023a], so we take both its architecture and its data pipeline as our starting point (Appendix G).

62 3 Method

63 **The base architecture.** DeepPeptide passes a frozen ESM-2 embedding through a CNN-BiLSTM
64 stack and decodes it with a linear-chain CRF that expands each of {Peptide, Propeptide} into a
65 chain of up to 50 position states, 101 in all, so a legal path can only realize a contiguous segment
66 of admissible length. The encoder computes one emission per label and shares it across all 50 states
67 of that label: the decoder is boundary-aware, the features feeding it are not. The two additions
68 close that gap from opposite ends of the pipeline (Figure 1b, drawn at reading size in Figure 2).
69 Other embedding sources and architectural modifications were screened as well, with verdicts in
70 Appendix F.

71 **Boundary head.** A small feed-forward block (LayerNorm, Linear, GELU, Linear, hidden size 64)
72 is applied to the contextual per-residue features. For each position and each segment type it emits
73 three numbers, scoring how much that position looks like a segment start, an interior position, or a
74 segment end, and these are added to the emissions of the corresponding position-fixed CRF states.
75 The final linear layer is initialized to zero, so training begins at exactly the base model and the head
76 only learns corrections that lower the loss. It costs 4,678 parameters on a trainable stack of 224,710.

77 **Adapter.** The second addition leaves the decoder alone and acts on the input. Before the per-
78 residue embedding reaches the CNN-BiLSTM it passes through LayerNorm, Linear, GELU, which
79 re-projects it and reduces its width (1280 to 256 for ESM-2). A pLM embedding is trained on
80 masked-residue recovery rather than cleavage-site localization, and its feature distribution suits nei-
81 ther this task nor the small head that consumes it, so a trainable re-projection in front of a frozen
82 encoder is the standard remedy. It is the more expensive of the two, a net +232,704 parameters.

83 4 Data and evaluation

84 4.1 Dataset

85 A precursor $x = (a_1, \dots, a_L)$ is labelled per residue with $y_t \in \{\text{None, Peptide, Propeptide}\}$
86 and contiguous runs of a label form typed segments, so the task is to recover which stretches of a
87 precursor become mature peptides and which become propeptides that are excised and discarded.
88 DeepPeptide built its dataset from PEPTIDE and PROPEP annotations in the 2022 Swiss-Prot release.
89 We rebuilt it with the same pipeline on the 2026 release [UniProt Consortium, 2025]. The collection
90 grows from 8,449 proteins to 9,619, of which 8,897 both carry ESM-2 embeddings and enter the
91 five folds used here. Folds come from GraphPart [Teufel et al., 2023a] at a 30% pairwise-identity
92 ceiling, balanced by cleavage-motif class. Segment-length filtering, motif balancing and the full
93 composition of the rebuild are given in Appendix E.

94 4.2 Evaluation criterion

95 Peptides and propeptides are matched separately. A prediction is a true positive when its start and its
96 end both fall within $\pm\tau$ residues of a true segment of the same type. Unmatched predictions are false
97 positives and unmatched true segments false negatives. Following the original evaluation we keep
98 $\tau = 3$ as the headline setting, since annotated sites carry a few residues of experimental uncertainty.
99 Sweeping τ down to zero then asks a second question, not whether a peptide was found but how
100 precisely its ends were placed. The reference implementation of this criterion carries an upstream
101 variable-shadowing bug that marks the wrong true segment as matched, so everything reported here
102 is re-scored with a corrected matcher (Appendix C).

103 **Agreement with the published baseline.** DeepPeptide reports precision 0.68 and recall 0.49 at
104 ± 3 , both segment types together, averaged over the twenty models of its nested cross-validation on
105 the 2022 data, which implies an F1 of 0.570. Running the base architecture on that release under the
106 same protocol scores 0.574 ± 0.069 with the original matcher and 0.590 ± 0.064 with the corrected
107 one of Appendix C. The F1 reproduces, the operating point does not: our precision is 0.614 and
108 our recall 0.541. The published system tunes its hyperparameters per outer fold, while we hold one
109 configuration fixed across all 20 cells so that every row of Table 1 is measured under the same setting

Table 1: 5×4 nested cross-validation with the corrected matcher. Each cell is the mean over the five outer folds, with the standard deviation across them as a subscript. Precision and recall are at the headline ± 3 tolerance, and the four F1 columns tighten it to an exact residue. *Growth* is how much a row’s F1 advantage over the ESM-2 base widens between ± 3 and an exact match.

Additions	at ± 3		F1 by tolerance				growth
	P	R	± 3	± 2	± 1	exact	
<i>ESM-2</i>							
none	0.621 $\pm$.025	0.539 $\pm$.045	0.576 $\pm$.029	0.544 $\pm$.027	0.467 $\pm$.020	0.354 $\pm$.010	—
boundary head	0.673 $\pm$.029	0.538 $\pm$.036	0.597 $\pm$.026	0.566 $\pm$.024	0.488 $\pm$.015	0.384 $\pm$.015	+0.010
adapter	0.628 $\pm$.019	0.579 $\pm$.040	0.602 $\pm$.026	0.572 $\pm$.024	0.491 $\pm$.011	0.396 $\pm$.020	+0.016
both	0.667 $\pm$.032	0.598 $\pm$.026	0.630 $\pm$.021	0.599 $\pm$.023	0.518 $\pm$.016	0.432 $\pm$.013	+0.023
<i>ESM-C 6B</i>							
none	0.620 $\pm$.017	0.560 $\pm$.027	0.588 $\pm$.016	0.558 $\pm$.013	0.483 $\pm$.007	0.371 $\pm$.013	+0.005
boundary head	0.731 $\pm$.010	0.572 $\pm$.030	0.641 $\pm$.022	0.608 $\pm$.026	0.532 $\pm$.035	0.429 $\pm$.030	+0.010
adapter	0.598 $\pm$.012	0.607 $\pm$.031	0.602 $\pm$.017	0.573 $\pm$.019	0.497 $\pm$.016	0.394 $\pm$.031	+0.014
both	0.690 $\pm$.013	0.646 $\pm$.033	0.666 $\pm$.018	0.636 $\pm$.015	0.554 $\pm$.016	0.466 $\pm$.021	+0.021

(Appendix D). On the 2026 rebuild the base reaches 0.576 ± 0.029 , so the rebuild does not move it either.

5 Results

We ran the experiments in two stages.

Screening. We first split the data into seven GraphPart folds with fixed roles: four for training, one for epoch selection, one for comparing architectures and one held back. Against that split we screened roughly a dozen modifications, with the protocol, the verdict table and the per-candidate figures in Appendix F.

The folds of that split are not interchangeable (Figure 4), and under one assignment of roles several modifications changed the *sign* of their measured effect between the two held-out folds. A single draw of this kind cannot resolve an effect of 0.02–0.03, so none of those numbers is reported as a finding. It can, however, separate the consistently unhelpful from the worth paying for, and two cleared that bar: the boundary head and the adapter sat at or above the base throughout.

Confirmation. We then re-evaluated the two survivors with the 5×4 nested cross-validation DeepPeptide itself uses, on ESM-2 and on ESM-C 6B, so that the embedding is a factor of the design and not a third candidate. For each of five outer folds we train four models on three of the remaining folds and validate on the fourth, then test on the outer fold, which they never saw. An outer fold’s score is the mean over its four models, the estimate is the mean over the five outer folds, and the reported spread is the standard deviation across outer folds rather than across all 20 cells, since four models sharing an outer fold share a test set. Because selection happened at the screening stage, what follows is a test of two pre-specified hypotheses.

Table 1 reports the same 2×2 factorial on both embeddings under that protocol, with the corrected matcher.¹

The two additions act on different errors. Each improves F1 by a similar amount on its own, and together they give 0.054 for ESM-2 and 0.079 for ESM-C, slightly more than the sum of the parts. Splitting F1 into its components explains why. The boundary head is almost purely a precision effect: precision rises by 0.052 while recall does not move at all (-0.001). The adapter is mostly a recall effect: recall rises by 0.040 against 0.007 of precision. Applied together they recover both, 0.046 precision and 0.059 recall.

¹Implementation, the run configuration behind every row and a synthetic smoke test: <https://anonymous.4open.science/r/cleavage-site-segmentation-5851>

139 **Tightening the tolerance.** The F1 columns of Table 1 sweep τ from ± 3 to an exact match. Ab-
140 solute F1 falls steeply for every model, the strongest configuration from 0.666 to 0.466. Placing a
141 cleavage site exactly is still an open problem.

142 A model that is better overall retains a different fraction of its ± 3 score even when no boundary is
143 placed better, so we compare absolute gaps rather than ratios. Every gap widens as the tolerance
144 tightens (the *growth* column of Table 1), and each stays nearly flat from ± 3 to ± 1 before opening
145 up at the last step, so what separates these models is specifically whether they hit the exact residue.

146 A widening gap is still not proof of better localization, since a model that finds more segments also
147 gains ground at every tolerance. To hold detection fixed we compared each variant against the base
148 on the true segments that *both* localize within ± 3 , matched cell by cell. On that set the share of
149 boundaries placed on the exact residue rises from 0.630 to 0.652 for the boundary head, from 0.619
150 to 0.668 for the adapter and from 0.621 to 0.698 for the two together, each positive on all five outer
151 folds. A gate is itself a choice, so Appendix B asks the same question with no window at all: when
152 a hit needs only a single residue of overlap, neither addition gains anything measurable (-0.013 for
153 the head, $+0.006$ for the adapter), and both gain steadily as the requirement tightens. What they buy
154 is placement rather than coverage, which is why a ± 3 window records part of it as recall. Scored at
155 individual cleavage sites, the head’s gain lies entirely on the C-terminal side.

156 **Capacity the base decoder cannot use.** The base architecture on ESM-C 6B reaches $0.588 \pm$
157 0.016 against 0.576 ± 0.029 on ESM-2, overlapping intervals, so an embedding twice as wide buys
158 no confirmed improvement on its own while either addition on the narrower one does. What the
159 additions are worth there depends on which one. Paired by outer fold against its own base, the
160 boundary head gains $+0.054 \pm 0.026$ on ESM-C 6B against $+0.021 \pm 0.007$ on ESM-2, the adapter
161 $+0.014 \pm 0.014$ against $+0.026 \pm 0.018$, and the two together $+0.079 \pm 0.009$ against $+0.054 \pm 0.016$,
162 all six positive on five folds of five. The adapter re-projects an embedding trained for masked-residue
163 recovery, and the less mismatched that embedding the less there is to re-project, whereas the head
164 supplies position-specific evidence and a richer embedding carries more of it. Read the other way,
165 the swap is worth $+0.012$ F1 under the base architecture and $+0.036$ under both additions, and
166 Appendix B reproduces the split on the protein-cells common to all runs. The capacity is in the
167 wider embedding either way. What changes is whether the decoder can reach it.

168 6 Conclusion

169 A zero-initialized boundary head at the decoder and a lightweight adapter at the input improve
170 segment F1 by $+0.054$ together on ESM-2 and by $+0.079$ on ESM-C 6B. They combine because
171 they act on different errors, the head suppressing spurious segments and the adapter recovering
172 true ones the base placed outside the window, and both place boundaries more precisely, so their
173 advantage widens as the tolerance tightens. The evaluation carries a second message: homology-
174 aware partitioning leaves the folds taxonomically non-exchangeable, and their spread exceeds either
175 addition measured on its own.

176 **The folds stay unequal.** Averaging over folds does not make them comparable, it only stops one
177 of them deciding the result. GraphPart keeps homologs together, and taxa *are* homology clusters: the
178 five outer folds differ in size by a factor of two (1,263 to 2,572 proteins), 313 of 714 *Conus* sequences
179 fall in fold 1 against 16 in fold 2, and all 293 *Cyriopagopus* avoid fold 0 entirely (Table 4). The base
180 architecture accordingly scores 0.530, 0.572, 0.576, 0.599 and 0.604 on the five outer folds, a range
181 of 0.074, three and a half times the $+0.021$ effect the same experiment resolves once averaged over
182 them.

183 **What the protocol does not settle.** Nothing was selected on outer-fold scores inside confirmation,
184 since every cell of the 2×2 is reported, but the screening that chose the candidates ran on the same
185 proteins the folds are drawn from, and sealing that would take a third held-out level. One cell costs
186 about 10 GPU-hours and the full $2 \times 2 \times 2$ grid roughly 1,600, so a third level multiplies that
187 again while cutting the training fraction from $3/5$ to $2/5$ of an already small dataset. At this size the
188 protocol cannot be made both clean and affordable. The spread we report covers fold composition
189 and not seed variation, so the true uncertainty is wider, and the base row is an untuned instance of
190 the architecture rather than the tuned published system (Appendix D).

References

- José Juan Almagro Armenteros, Marco Salvatore, Olof Emanuelsson, Ole Winther, Gunnar von Heijne, Arne Elofsson, and Henrik Nielsen. Detecting sequence signals in targeting peptides using deep learning. *Life Science Alliance*, 2(5):e201900429, 2019. doi: 10.26508/lsa.201900429.
- Zhenjiao Du, Xingjian Ding, William Hsu, Arslan Munir, Yixiang Xu, and Yonghui Li. pLM4ACE: a protein language model based predictor for antihypertensive peptide screening. *Food Chemistry*, 431:137162, 2024. doi: 10.1016/j.foodchem.2023.137162.
- Peter Duckert, Søren Brunak, and Nikolaj Blom. Prediction of proprotein convertase cleavage sites. *Protein Engineering, Design and Selection*, 17(1):107–112, 2004. doi: 10.1093/protein/gzh013.
- Ahmed Elnaggar, Hazem Essam, Wafaa Salah-Eldin, Walid Moustafa, Mohamed Elkerdawy, Charlotte Rochereau, and Burkhard Rost. Ankh: Optimized protein language model unlocks general-purpose modelling, 2023. URL <https://arxiv.org/abs/2301.06568>.
- ESM Team. Esm cambrian: Revealing the mysteries of proteins with unsupervised learning. <https://www.evolutionaryscale.ai/blog/esm-cambrian>, December 2024.
- Thomas Hayes, Roshan Rao, Halil Akin, Nicholas J. Sofroniew, Deniz Oktay, Zeming Lin, Robert Verkuil, Vincent Q. Tran, Jonathan Deaton, Marius Wiggert, Rohil Badkundri, Irhum Shafkat, Jun Gong, Alexander Derry, Raul S. Molina, Neil Thomas, Yousuf A. Khan, Chetan Mishra, Carolyn Kim, Liam J. Bartie, Matthew Nemeth, Patrick D. Hsu, Tom Sercu, Salvatore Candido, and Alexander Rives. Simulating 500 million years of evolution with a language model. *Science*, 387(6736):850–858, 2025. doi: 10.1126/science.ads0018.
- Michael Heinzinger, Konstantin Weissenow, Joaquin Gomez Sanchez, Adrian Henkel, Milot Mirdita, Martin Steinegger, and Burkhard Rost. ProstT5: bilingual language model for protein sequence and structure. *NAR Genomics and Bioinformatics*, 6(4):lqae150, 2024. doi: 10.1093/nargab/lqae150.
- Fuyi Li, Jinxiang Chen, André Leier, Tatiana Marquez-Lago, Quanzhong Liu, Yanze Wang, Jerico Revote, A. Ian Smith, Tatsuya Akutsu, Geoffrey I. Webb, Lukasz Kurgan, and Jiangning Song. DeepCleave: a deep learning predictor for caspase and matrix metalloprotease substrates and cleavage sites. *Bioinformatics*, 36(4):1057–1065, 2020a. doi: 10.1093/bioinformatics/btz721.
- Fuyi Li, André Leier, Quanzhong Liu, Yanan Wang, Dongxu Xiang, Tatsuya Akutsu, et al. Procleave: predicting protease-specific substrate cleavage sites by combining sequence and structural information. *Genomics, Proteomics & Bioinformatics*, 18(1):52–64, 2020b. doi: 10.1016/j.gpb.2019.08.002.
- Fuyi Li, Cong Wang, Xudong Guo, Tatsuya Akutsu, Geoffrey I. Webb, Lachlan J. M. Coin, Lukasz Kurgan, and Jiangning Song. ProsperousPlus: a one-stop and comprehensive platform for accurate protease-specific substrate cleavage prediction and machine-learning model construction. *Briefings in Bioinformatics*, 24(6):bbad372, 2023. doi: 10.1093/bib/bbad372.
- Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023. doi: 10.1126/science.ade2574.
- Catherine Mooney, Niall J. Haslam, Gianluca Pollastri, and Denis C. Shields. Towards the improved discovery and design of functional peptides: common features of diverse classes permit generalized prediction of bioactivity. *PLoS ONE*, 7(10):e45012, 2012. doi: 10.1371/journal.pone.0045012.
- Catherine Mooney, Niall J. Haslam, Thérèse A. Holton, Gianluca Pollastri, and Denis C. Shields. PeptideLocator: prediction of bioactive peptides in protein sequences. *Bioinformatics*, 29(9):1120–1126, 2013. doi: 10.1093/bioinformatics/btt103.

- Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo, Myle Ott, C. Lawrence Zitnick, Jerry Ma, and Rob Fergus. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences*, 118(15):e2016239118, 2021.
- Jiangning Song, Hao Tan, Andrew J. Perry, Tatsuya Akutsu, Geoffrey I. Webb, James C. Whisstock, and Robert N. Pike. PROSPER: an integrated feature-based tool for predicting protease substrate cleavage sites. *PLoS ONE*, 7(11):e50300, 2012. doi: 10.1371/journal.pone.0050300.
- Jiangning Song et al. PROSPERous: high-throughput prediction of substrate cleavage sites for 90 proteases with improved accuracy. *Bioinformatics*, 34(4):684–687, 2018. doi: 10.1093/bioinformatics/btx670.
- Bruce R. Southey, Andinet Amare, Tyler A. Zimmerman, Sandra L. Rodriguez-Zas, and Jonathan V. Sweedler. NeuroPred: a tool to predict cleavage sites in neuropeptide precursors and provide the masses of the resulting peptides. *Nucleic Acids Research*, 34(Web Server issue):W267–W272, 2006. doi: 10.1093/nar/gkl161.
- Felix Teufel, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Magnús Halldór Gíslason, Silas Irby Pihl, Konstantinos D. Tsirigos, Ole Winther, Søren Brunak, Gunnar von Heijne, and Henrik Nielsen. SignalP 6.0 predicts all five types of signal peptides using protein language models. *Nature Biotechnology*, 40(7):1023–1025, 2022. doi: 10.1038/s41587-021-01156-3.
- Felix Teufel, Magnús Halldór Gíslason, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Ole Winther, and Henrik Nielsen. GraphPart: homology partitioning for biological sequence analysis. *NAR Genomics and Bioinformatics*, 5(4):lqad088, 2023a. doi: 10.1093/nargab/lqad088.
- Felix Teufel, Jan Christian Refsgaard, Christian Toft Madsen, Carsten Stahlhut, Mads Grønberg, Ole Winther, and Dennis Madsen. Deeppeptide predicts cleaved peptides in proteins using conditional random fields. *Bioinformatics*, 39(10):btad616, 10 2023b. ISSN 1367-4811. doi: 10.1093/bioinformatics/btad616. URL <https://doi.org/10.1093/bioinformatics/btad616>.
- UniProt Consortium. UniProt: the universal protein knowledgebase in 2025. *Nucleic Acids Research*, 53(D1):D609–D617, 2025. doi: 10.1093/nar/gkae1010. URL <https://academic.oup.com/nar/article/53/D1/D609/7902999>.
- Michel van Kempen, Stephanie S. Kim, Charlotte Tumescheit, Milot Mirdita, Jeongjae Lee, Cameron L. M. Gilchrist, Johannes Söding, and Martin Steinegger. Fast and accurate protein structure search with Foldseek. *Nature Biotechnology*, 42:243–246, 2024. doi: 10.1038/s41587-023-01773-0.
- Lei Wang, Chen Huang, Mingxia Wang, Zhidong Xue, and Yan Wang. NeuroPred-PLM: an interpretable and robust model for neuropeptide prediction by protein language model. *Briefings in Bioinformatics*, 24(2):bbad077, 2023. doi: 10.1093/bib/bbad077.
- Lei Wang, Zilu Zeng, Zhidong Xue, and Yan Wang. DeepNeuropePred: a robust and universal tool to predict cleavage sites from neuropeptide precursors by protein language model. *Computational and Structural Biotechnology Journal*, 23:309–315, 2024. doi: 10.1016/j.csbj.2023.12.004.
- M. Zhang, J. Zhou, X. Wang, et al. DeepBP: ensemble deep learning strategy for bioactive peptide prediction. *BMC Bioinformatics*, 25:352, 2024. doi: 10.1186/s12859-024-05974-5.
- Lei Zhu, Hao Sun, and Sen Yang. BPFun: a deep learning framework for bioactive peptide function prediction using multi-label strategy by transformer-driven and sequence rich intrinsic information. *BMC Bioinformatics*, 26(1):187, 2025. doi: 10.1186/s12859-025-06190-5.

Impact Statement

This paper improves a computational tool for predicting proteolytic cleavage products from protein sequence. Potential applications include interpreting proteomics data, informing vaccine and peptide-drug design, and studying disease-associated proteolysis. We are not aware of societal risks specific to this work beyond those generally associated with more accurate protein sequence analysis tools.

289 Appendix

290 A The architecture at reading size




Figure 2: The base architecture and the two additions, at reading size. Figure 1b shows the same diagram at the scale of a single panel, where it locates the two blocks rather than documenting them. Top: the pipeline, with the adapter between the frozen pLM and the CNN-BiLSTM encoder and the boundary head reading the encoder’s output. Bottom: what the head adds. The base model computes one emission per label and shares it across all 50 position states of that label, so the CRF’s start, interior and end states are fed identical evidence. The head emits three numbers per position and per segment type and adds them to the states they describe.

291 B Boundary placement beyond the ± 3 window

292 Section 5 makes two claims that no single tolerance can settle: that the additions place boundaries
 293 better rather than merely finding more segments, and that what they buy survives a criterion with no
 294 acceptance window at all. This appendix reports the evidence for both.

295 **Holding detection fixed.** The comparison in Section 5 is paired. For each cell we take the true
 296 segments that the variant and the base both localize within ± 3 , so both models have found the
 297 segment and only its placement is at stake, and ask how often each puts a boundary on the annotated
 298 residue. Pooled over the twenty cells of a configuration this leaves 27,013 segments for the boundary
 299 head, 29,277 for the adapter and 29,394 for the two together.

300 **What each block does to the count.** Those denominators are themselves the result. Of the true
 301 segments the base localizes within ± 3 across the twenty cells, the boundary head localizes 320 fewer
 302 and the adapter 2,236 more, with both together 3,289 more. The head therefore improves placement
 303 while covering marginally less, which is the precision effect of Table 1 seen segment by segment,
 304 and the adapter does the opposite. On the paired set the head moves 13.2% of boundaries closer to
 305 the annotated residue and 11.4% further away, a near-even trade that nets out positive, while the two
 306 together move 15.0% closer against 7.5% further.

307 **Dropping the window.** A gate is itself a choice, so we also score the same predictions with no
 308 tolerance. Predicted and true segments of the same type are matched greedily one to one by inter-
 309 section over union, and a match counts as a hit when the overlap clears a threshold τ . Precision
 310 and recall use all predicted and all true segments as their denominators, so nothing is conditioned
 311 on what either model happened to find. Every configuration is scored on the 35,504 protein-cells
 312 common to all runs, which takes the coverage difference between the two embeddings (8,897 against
 313 8,999 proteins) out of the comparison. On that common set the ESM-2 to ESM-C 6B swap is worth

314 +0.005 F1 at $\text{IoU} \geq 0.5$ under the base architecture, on two outer folds of five, and +0.024 under
 315 both additions, on five of five, reproducing the split that Section 5 measures at ± 3 .

Table 2: Segment F1 under overlap matching. The first column requires only that a predicted and a true segment of the same type share one residue, and the rest require an intersection over union of at least τ . Same cells as Table 1, restricted to the protein set common to every run, mean over the five outer folds.

	overlap	$\tau \geq 0.5$	$\tau \geq 0.8$	$\tau \geq 0.9$	$\tau \geq 0.95$
<i>ESM-2</i>					
base	0.729	0.667	0.569	0.487	0.411
+ boundary head	0.715	0.668	0.585	0.515	0.438
+ adapter	0.734	0.686	0.590	0.516	0.441
+ both	0.746	0.703	0.616	0.546	0.474
<i>ESM-C 6B</i>					
base	0.738	0.672	0.574	0.496	0.419
+ boundary head	0.735	0.704	0.624	0.555	0.478
+ adapter	0.729	0.681	0.587	0.514	0.440
+ both	0.759	0.727	0.649	0.582	0.507

316 **Placement, not coverage.** The looser the criterion, the less either addition is worth. Paired by
 317 outer fold against the ESM-2 base, at a gate that asks only for a single residue of overlap the boundary
 318 head is worth -0.013 ± 0.026 and the adapter $+0.006 \pm 0.012$, neither distinguishable from zero.
 319 At $\tau \geq 0.5$ they are worth $+0.001 \pm 0.021$ and $+0.019 \pm 0.017$, and at $\tau \geq 0.9$ $+0.028 \pm 0.004$ and
 320 $+0.028 \pm 0.015$, both on five folds of five. Underneath that near-zero F1 at the overlap gate the two
 321 move in opposite directions: the head issues 219 fewer segments per fold, buying $+0.022$ precision
 322 for -0.035 recall, while the adapter issues 130 more and buys $+0.024$ recall for -0.021 precision.
 323 Neither finds appreciably more of the annotation than the base architecture does. What changes is
 324 where the boundaries land. This is the same conclusion Table 1 reaches through a tolerance, reached
 325 without one.

326 The mean IoU taken over *all* true segments moves the other way for the head, 0.576 to 0.559, while
 327 its F1 rises at every threshold. That is the precision effect of Table 1 seen from another angle: the
 328 head withdraws weak predictions, so some true segments lose a poor match altogether, and the ones
 329 that keep a match are placed better.

330 **Which end moves.** Scoring cleavage sites rather than segments, at exact placement and against
 331 all annotated sites, separates the two additions again. The boundary head improves the C-side of a
 332 segment, $+0.028 \pm 0.006$ on five folds of five, and leaves the N-side alone, -0.007 ± 0.020 on two of
 333 five. The adapter improves both, $+0.022 \pm 0.006$ and $+0.028 \pm 0.009$, and the two together improve
 334 both further, $+0.042 \pm 0.010$ and $+0.060 \pm 0.006$. Split by segment type, the head’s C-side gain
 335 is largest on propeptides ($+0.037 \pm 0.021$, five of five), whose C-terminal cut is where the mature
 336 peptide begins and is the weaker of the two C-sides: the base model scores 0.476 there against 0.525
 337 on peptide C-sides.

338 C Metric implementation bug

339 While auditing the DeepPeptide implementation of the matching criterion (Section 4.2), we found
 340 a bug in the function that implements it (`get_counts_for_protein`), reproduced below in simpli-
 341 fied form:

```

342 for idx, row in true_df.iterrows(): # idx: true-segment index
343     true_start, true_stop = row['start'], row['stop']
344     for idx, row in pred_df.iterrows(): # idx overwritten with pred index
345         pred_start, pred_stop = row['start'], row['stop']
346         ...
347         if start_match and stop_match:
348             true_df.loc[idx, 'matched'] = True # BUG: idx is now the
349             pred_df.loc[idx, 'matched'] = True # predicted index, not the true one
350             break
351 
```

The outer loop iterates over true segments and the inner loop over predicted segments, but both reuse the loop variable `idx`. By the time a match is found, `idx` refers to the predicted segment, so `true_df.loc[idx, 'matched']` marks the wrong true segment. The effect is not large on average (it is partly self-correcting), but it costs a true positive whenever it fires, so both recall and precision come out low, and F1 with them. Across the nested-CV grid the correction raises recall by 0.021 to 0.032 and precision by 0.005 to 0.010. The loss concentrates in proteins where the model predicts more segments than the protein actually has. This bug is present in the upstream DeepPeptide code, not something we introduced.

We did not silently patch this function in place, since that would shift absolute numbers without documenting it. Instead we implemented a separate, corrected matcher and used it to re-score the results in this paper: the earlier single-split runs reported in Appendix F, except two whose model classes no longer exist in the code and whose checkpoints cannot be rebuilt, and the full 5×4 nested-cross-validation grid in Table 1 (all 160 cells of the eight configurations, and the 20 further cells of the 2022-release reproduction), by re-running test-partition inference from each cell’s saved checkpoint on the machine that held them. Every table in the main text uses the corrected matcher, and the original numbers are kept alongside in the released run artifacts.

Re-scoring the grid also recomputed the original matcher on the same decoded segments, so each cell could be checked against the number already recorded for it. Over the 80 cells for which that check could be run, the two agree to 1.2×10^{-3} F1 at worst and to 3×10^{-7} at the median, so what remains is inference nondeterminism rather than disagreement between two deterministic functions of the same segments.

Across the seven runs whose every cell carries both matchers the correction shifts mean F1 by +0.015 to +0.022, always upward, since the bug was discarding true positives. On ESM-2 it leaves the ordering of the four configurations intact, base < boundary head < adapter < combined, before and after correction. That ordering is not robust and we build no argument on it: the boundary-head/adaptor gap is 0.005 against a fold-level standard deviation of 0.026, and per outer fold the full ordering holds in four folds of five after correction and three of five before it. On ESM-C 6B the ordering differs, base < adapter < boundary head < combined, the two middle configurations swapping places.

D Relation to the published DeepPeptide system

Our base architecture is DeepPeptide’s architecture, but it is not DeepPeptide’s published model, and the difference is worth stating because it is the reason Section 4.2 reproduces the published F1 without reproducing the published precision and recall.

The released implementation spends the inner loop of its nested cross-validation on a hyperparameter search, and the twenty checkpoints behind the published numbers carry a different winning configuration for each outer fold: learning rate from 3.3×10^{-4} to 5.5×10^{-3} , dropout from 0.03 to 0.69, convolutional dropout from 0.04 to 0.51, kernel size 3 or 5, 48 to 96 filters, hidden size 32 or 48, and batch size 20 to 90. We hold a single configuration fixed across all twenty cells at the implementation’s defaults (learning rate 10^{-4} , dropout and convolutional dropout 0.1, kernel size 3, 32 filters, hidden size 64, batch size 48) and spend the inner loop on epoch selection alone.

That choice is deliberate. A factorial is only interpretable if its cells differ in the factor under study and in nothing else, and per-fold tuning would confound each addition with whatever the search happened to find for it. The cost is that our base row is an untuned instance of the architecture rather than the tuned system, so the gains in Table 1 are gains over that instance. They should not be read as gains a per-fold hyperparameter search would leave on the table, and every configuration in the grid carries the same handicap.

Two further notes on the reference implementation. Its recall guard tests the wrong denominator, returning zero when no segment is predicted but dividing by the true-segment count that the guard never checked; on this data the branch is unreachable, so the value is latent rather than wrong. And its evaluation script builds the ground truth with overlapping annotations kept separate while its training dataset merges them, a discrepancy we did not attempt to reconcile because our scoring rebuilds the ground truth from the labelled sequences directly.

405 E Dataset details

Table 3: Dataset composition: the 2022 release used by the original DeepPeptide against the 2026 release rebuilt here. Counts are before homology partitioning.

	2022	2026
Proteins	8449	9619
Peptide segments	6372	7431
Propeptide segments	8211	9140
Median peptide length (residues)	21	20




Figure 3: Composition of the rebuilt 2026 dataset: segment length, type, and genus distributions.

406 Two filters from the original pipeline are kept unchanged. Segments shorter than 5 or longer than
 407 50 residues are dropped: the CRF’s state count fixes the upper bound, and a segment shorter than
 408 five residues is barely scoreable at a ± 3 tolerance. Folds are then balanced by cleavage-motif class,
 409 obtained by k -means ($k = 50$) on ESM-2 embeddings of the four residues flanking each annotated
 410 boundary, on top of the 30% pairwise-identity ceiling GraphPart enforces between folds.

Table 4: Proteins per outer fold, and how three frequent genera distribute across them. *Cyriopagopus* and *Lycosa* are spider venoms, *Conus* a cone snail, and each is a homology cluster that GraphPart is obliged to keep intact. *Homo*, by contrast, is spread evenly (71/81/66/130/84).

	total	f0	f1	f2	f3	f4
All proteins	8897	1558	2572	1263	2025	1479
<i>Conus</i>	714	160	313	16	149	76
<i>Cyriopagopus</i>	293	0	62	45	9	177
<i>Lycosa</i>	163	5	42	10	106	0

411 F Screening protocol and verdicts

412 F.1 Protocol

413 Data was split into seven GraphPart folds (30% identity threshold, motif-balanced as in Section 4.1),
 414 with folds assigned fixed roles: four folds for training, one for validation (epoch selection), one

415 for model selection (comparing architectures), and one originally intended as a fully sealed test
 416 fold. This separated architecture selection from evaluation, but the evaluation itself was still a single
 417 draw: one random assignment of roles, one pair of held-out folds. Several modifications changed
 418 the sign of their measured effect between the two held-out folds. Replacing ESM-2 with ESM-C 6B
 419 measured $+0.074$ on one and -0.011 on the other, ESM-C 6B against ESM-C 600M gave $+0.028$
 420 and -0.021 , and the net effect of the 3Di channel gave $+0.015$ and -0.018 . Effects that kept their
 421 sign still moved by a lot: the isolated gated adapter measured -0.001 on one fold and $+0.045$ on
 422 the other. Figure 4 shows one reason, namely that the folds differ substantially in segment-length
 423 composition. The two held-out folds turn out to be complementary in segment-length composition,
 424 fold 2 concentrated at 10–24 residues and fold 5 bimodal, so results below pool them (model-select
 425 $\cup$ sealed, $\approx 2,300$ proteins) to reduce (but not eliminate) this instability. This pooling is why the
 426 “sealed” fold is not, in practice, a held-out test set independent of model selection.




Figure 4: Why the folds of the earlier seven-fold split are not interchangeable: (a) the segment-length profile of each fold against the profile of the whole dataset, (b) the L_1 distance between the two. Fold 4 tracks the overall distribution closely (0.07) while five of the seven folds sit above 0.23. Balancing was done on cleavage motif and homology, not on segment length, so this axis was left free. The same holds for the five-fold split used in the main text, where the imbalance is taxonomic (Table 4).

427 **F.2 Summary of tested modifications**

428 Table 5 summarizes the modifications tested under screening protocol, with the corrected matcher
 429 (Appendix C) applied throughout. Two of these rows were carried forward as candidates, the bound-
 430 ary head and the adapter (Section 3), and both were confirmed. The embedding swap was not a
 431 candidate but became a factor of the confirmation design, and it is the row this table reads least well:




Figure 5: Precision/recall trade-offs by segment type when adding the 3Di channel or the bond-prediction loss, single-split protocol, pooled held-out folds.

432 nested cross-validation puts ESM-C 6B at 0.588 ± 0.016 against ESM-2's 0.576 ± 0.029 (Table 1),
 433 overlapping intervals, where this table records a confident $+0.03$. Everything else below remains at
 434 the confidence of a single pooled split.

Table 5: Modifications tested under the screening single-split protocol, with their measured effect on F1 (pooled held-out folds) and an informal verdict. None of these effect sizes should be compared directly to the nested-CV numbers in Table 1.

Modification	Type	Measured effect on F1	Verdict
ESM-C 6B instead of ESM-2	embedding swap	$+0.03$, unstable across folds	helps
Boundary head on ESM-C 6B	decoder addition	$+0.05$, $+0.07$, consistent	helps
Boundary head on ESM-2	decoder addition	≈ 0 (CI includes zero)	no effect
Structural channel (3Di)	extra input	propep. $+0.02$ / pep. -0.03 , net trade-off	no effect
ESM-C 6B compression 2560 $\rightarrow$ 256	optimization	≈ 0 , $10\times$ narrower input and $16\times$ fewer trainable parameters	no effect
Gated adapter on ESM-C 6B (pLM re-projection)	input adapter	$+0.022$, CI $[+0.008, +0.038]$	helps
Bond-prediction auxiliary loss (on ESM-C 6B)	extra loss	pep. -0.05 , net -0.015 , neutral on ESM-2	harmful
Telescopic segment CRF	decoder addition	≈ 0 on ESM-2, $+0.022$ on ESM-C 6B, no CI available	unresolved

435 Note the apparent tension with Table 1: under this earlier protocol, the boundary head on ESM-2
 436 looked like it had no effect, while nested cross-validation resolved a confirmed $+0.021$ F1 gain
 437 for the same modification on the same embedding. We read this as evidence that the single-split
 438 protocol lacked the statistical power to resolve an effect of this size, not as a contradiction – and as
 439 the clearest illustration of why we do not treat single-split numbers as findings in this paper. Whether
 440 the boundary head's effect is specifically larger on richer embeddings (as the ESM-C 6B rows above
 441 suggest) is accordingly left as an open question, not stated as a conclusion here.

F.3 Additional exploratory modifications

A larger set of ideas was tried under the same single-split protocol and did not show a clear, consistent benefit. We list them for completeness and as candidates for future re-testing, without effect-size claims:

- **Known-peptide dictionary.** An Aho–Corasick index of peptides from the training data was used to flag substring matches in a query sequence, injected as a bonus on the CRF state emissions, as a bias on the start, inside and end emissions specifically, fused into the encoder hidden state, and concatenated with the embedding before the encoder. No consistent benefit was observed, and the checkpoint for one of the variants can no longer be loaded.
- **Structural features.** Features extracted from predicted 3D structures (AlphaFold2 representations via the AFToolkit codebase, with several confidence-based filtering variants) and from the ProstT5 [Heinzinger et al., 2024] structural alphabet (3Di, as used in Foldseek [van Kempen et al., 2024]) were tested as additional input channels, and results were mixed (Table 5).
- **LoRA fine-tuning.** Low-rank adaptation of the last few pLM layers, instead of fully frozen embeddings, scored 0.558 and 0.567 against 0.621 for the frozen baseline, all three on the 2022 release under its own five-fold split rather than the protocol of Appendix F. It also ran for fewer epochs in the same time budget, though we did not measure the cost directly.
- **Projector variants.** Multi-scale and multi-branch projectors between the embedding and the CNN–BiLSTM were tried as alternative re-projection schemes to the adapter of Section 3.
- **Structural-projection width and telescopic CRF.** The width of the structural-feature projection (16/32/48 units) left pooled F1 flat at 0.691, 0.697 and 0.693, the telescopic segment CRF, an alternative decoder formulation that scores whole segments through a relative-position head, came out level with the base decoder on ESM-2 and +0.022 ahead on ESM-C 6B. Per-protein outputs were not retained for either run, so neither number carries a confidence interval and we treat the modification as untested rather than as having no effect.

None of these modifications showed a consistent gain large enough, under the single-split protocol, to justify testing under nested cross-validation ahead of the two reported in the main text. The complete experiment log (including runs not summarized above) is maintained in the project repository rather than reproduced here, since it was collected under a protocol we no longer treat as sufficient evidence on its own.

F.4 What the screening runs looked like

Three figures from the screening stage are worth keeping, each for a different reason, and none of them should be read as confirmed evidence in the sense of Section 5. Figure 7 is the comparison the verdict column of Table 5 summarizes, with the bootstrap intervals that made most of those verdicts unsafe. Figure 8 is the earliest form of the interaction that Section 5 later confirms under nested cross-validation, and shows that the boundary head was already worth more on the wider embedding a protocol ago. Figure 9 asks what recall actually tracks: for both ESM-C models it follows the maximum sequence identity between a held-out peptide and any training segment far more closely than it follows how many training segments the protein’s genus contributes, which is the behaviour of a model retrieving near neighbours rather than one that has learned the cleavage grammar. On the ESM-2 baseline the two axes are closer to comparable. This is a single split and we draw no conclusion from it, but it is the observation that would most repay a proper test.




Figure 6: F1 (± 3) against the number of training proteins, single-split protocol. ESM-2 rises from 0.498 at 40% of the training folds to 0.583 at 85%, then falls back to 0.572 at 100%, so it is still gaining over most of the range but not at the last point. The ESM-C 6B curves are flat from the smallest size tested, with movement smaller than the bootstrap interval, so nothing is observed saturating: they simply never climb.




Figure 7: Screening comparison on the pooled held-out folds at ± 3 : F1, precision and recall with bootstrap confidence intervals over 2,322 proteins. The top row is a gated-projector control at full width, not the configuration carried into the main text.




Figure 8: Boundary head × embedding interaction under the screening protocol: F1 gain from adding the boundary head on top of ESM-2 against on top of ESM-C 6B. This is the earlier, unconfirmed version of the interaction that Section 5 measures under nested cross-validation.




Figure 9: Recall on held-out peptides as a function of (a) how well-represented the protein’s genus is in training and (b) maximum sequence identity to a training segment, screening protocol. For the two ESM-C models recall tracks maximum identity to a training segment far more than genus abundance, while for the ESM-2 baseline the two axes are closer to comparable. The x axis in (a) counts training *segments* in the genus, not proteins.

G Extended review of prior work

Methods for predicting proteolytic peptides fall into three broad classes.

Protease-specific models. These predict cleavage only for a limited set of enzymes, most of which have a known recognition motif. ProP [Duckert et al., 2004] targets cleavage by the PACE/PC family in animals and plants. PROSPER [Song et al., 2012] and its successor PROSPERous [Song et al., 2018] predict cleavage for one chosen enzyme at a time, as does the later ProsperousPlus [Li et al., 2023]. DeepCleave [Li et al., 2020a] is restricted to caspases and matrix metalloproteinases, and Procleave [Li et al., 2020b] likewise handles one enzyme at a time and additionally requires 3D structural input. The shared limitation of this class is scope: in practice, one is usually interested in the combined effect of all proteases active in an organism, not any single enzyme.

Organism- or peptide-type-specific models. A separate class targets neuropeptides or other bioactive fragments specific to a given organism. NeuroPred [Southey et al., 2006] and the species-agnostic DeepNeuropePred [Wang et al., 2024] only detect cleavage near a handful of basic residues, while NeuroPred-PLM [Wang et al., 2023] does not localize cleavage sites at all but classifies an already-excised sequence. A related group of models finds signal subsequences rather than any bioactive peptide: SignalP [Teufel et al., 2022] and TargetP [Almagro Armenteros et al., 2019] identify the type of signal peptide, which indicates where a protein is trafficked.

General-purpose models on pLM embeddings. Protein language models (pLMs) are transformers trained without supervision on large sequence corpora, typically by masked-residue recovery, and they produce transferable per-residue embeddings: ESM-2 [Lin et al., 2023], the earlier ESM line [Rives et al., 2021], ESM Cambrian [ESM Team, 2024, Hayes et al., 2025], and Ankh [El-naggar et al., 2023]. Compact task-specific heads are then trained on top of these frozen embeddings. Besides **DeepPeptide** [Teufel et al., 2023b] itself, this recipe underlies DeepNeuropePred, NeuroPred-PLM, and SignalP 6, as well as models that score already-excised peptides rather than finding them: pLM4ACE [Du et al., 2024] predicts ACE-inhibitory activity, and BPFun [Zhu et al., 2025] and DeepBP [Zhang et al., 2024] predict broader bioactivity, building on earlier pre-pLM work such as PeptideRanker [Mooney et al., 2012]. Before pLMs, PeptideLocator [Mooney et al., 2013] addressed a problem close to ours, but its output is a per-residue heatmap of similarity to known bioactive peptides rather than an explicit segmentation, and it was the main point of comparison in the original DeepPeptide paper.

Baseline architecture: DeepPeptide. Since we build directly on it throughout this paper (Section 3), DeepPeptide’s architecture is worth describing in detail rather than treating as a black box. It first encodes a precursor sequence with a frozen pLM (ESM-2), producing a per-residue embedding that is not fine-tuned during training. This embedding is refined by a convolution, a single-layer bidirectional LSTM and a second convolution. The published description does not fix the widths: they are searched with Optuna inside the inner cross-validation loop. The configuration we train throughout uses 32 convolutional filters of width 3 and an LSTM hidden size of 64 per direction, giving 224,710 trainable parameters for the base model. These features are consumed by a linear-chain CRF decoder whose state set is extended well beyond the three raw labels {None, Peptide, Propeptide}: each of the two positive labels is expanded into a chain of states long enough to represent segments up to a fixed maximum length (101 states in total for the 50-residue cap used here, Section 4.1), so that a legal path through the CRF can only realize a contiguous segment of valid length, with dedicated states marking its first and last positions. This makes segmentation, rather than independent per-residue classification, the object the model is directly optimized for – and it is also where the gap described in Section 1 lives: every state in this extended set that shares a type receives the same emission from the encoder, regardless of whether it marks the start, interior, or end of a segment.

Baseline dataset for this task. Beyond its architecture, DeepPeptide also established the data resource this line of work relies on: precursor sequences from UniProtKB/Swiss-Prot annotated with PEPTIDE and PROPEP feature types, partitioned into homology-aware folds with GraphPart [Teufel et al., 2023a] at a 30% pairwise-identity threshold and additionally balanced by cleavage-flanking motif, then evaluated with the same family of nested cross-validation we adopt (Section 5). This gave the field a standardized, already homology-controlled benchmark rather than an ad hoc collection of previously published peptide lists, and both the curated data and the construction pipeline that produced it are public. Section 4.1 describes what we do with that resource: rebuilding it on a

542 current UniProtKB/Swiss-Prot release, which adds 1,170 proteins net and refreshes the annotations
543 behind every segment.

544 Among these, DeepPeptide is, to our knowledge, the only model that poses the task as full segmen-
545 tation of the precursor into typed peptide and propeptide segments, and it is the strongest available
546 baseline for this formulation. We therefore use its architecture and its dataset-construction method-
547 ology as the starting point for this work, rather than treating it as a system to audit: our contribution
548 is an architectural addition to this general recipe (Section 3), evaluated under the same 5×4 nested
549 cross-validation it uses (Section 4) on data rebuilt from a current release (Section 4.1).