
How Sharp Is a Cleavage Oracle? Position-Aware CRF Emissions, an Embedding Adapter, and What a ± 3 Benchmark Hides

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Designing a peptide construct means predicting where proteases will cut it, and
2 a design loop can only use a cleavage predictor at the resolution that predictor
3 is measured at. The benchmark for this task accepts a predicted segment when
4 both of its endpoints land within ± 3 residues of the annotation, which on a 5–50
5 residue peptide is most of the peptide: an oracle can improve its score by finding
6 more segments while placing none of them better, and can place every bound-
7 ary exactly and gain almost nothing. We make placement the target of both the
8 model and the measurement. Two small additions to the strongest available seg-
9 menter, a zero-initialized boundary head that gives its conditional random field
10 the position-specific evidence its own state graph already distinguishes, and a
11 lightweight adapter that re-projects the frozen protein language model embedding,
12 improve segment F1 by +0.054 together on ESM-2 and +0.079 on ESM-C 6B
13 under the 5×4 nested cross-validation the benchmark itself uses. Scored without
14 any acceptance window, the head covers less of the annotation than the base and
15 the adapter covers more, so most of what the pair buys is placement rather than
16 coverage, which is what a design loop needs and what a ± 3 window cannot report.
17 We also quantify why this benchmark is hard to measure on at all: homology-
18 aware folds are taxonomically non-exchangeable, and their spread exceeds either
19 addition measured on its own.

20 1 Introduction

21 Peptide and protein therapeutics are designed as precursors, and what reaches the target is whatever
22 the proteases of the host leave behind. A multi-epitope vaccine construct is the clearest case: the
23 fragments it presents are cut out in vivo, so the construct has to be checked at the design stage for the
24 peptides it will actually yield and for junctions prone to unwanted cleavage. A model that predicts
25 those cuts is an oracle in the design loop, and the loop can only use it at the resolution at which the
26 oracle has been measured.

27 That resolution is currently coarse, and the coarseness is in the benchmark rather than in the models.
28 A predicted segment counts as correct when both of its endpoints fall within ± 3 residues of a ground-
29 truth segment of the same type, and a single F1 at that tolerance is reported (Figure 1). Segments
30 are 5–50 residues long, so on a short peptide the acceptance window is comparable to the peptide
31 itself. Two consequences follow, and they pull in opposite directions. An oracle can raise its score
32 by finding more segments while placing none of them better. And it can place every boundary on the
33 annotated residue and gain almost nothing, because predictions that close were already counted. For
34 a design loop the second is the expensive one: moving a cut by three residues changes the peptide,
35 and the metric that ranks oracles cannot see the difference.




Figure 1: The problem, the two additions and what they buy. (a) One held-out precursor, with the ± 3 acceptance window shaded around each annotated cleavage site. (b) The architecture, with the two additions marked on it. (c) The F1 gap to the ESM-2 base at each tolerance, the base itself drawn flat at zero as the reference every other line is measured against. A curve is flat only if it is a parallel translation of the baseline, and none of the six carrying an addition is. Corrected matcher throughout, means over the five outer folds.

Two of our three contributions are accordingly about measurement: a decomposition of what each block buys, established at a criterion with no acceptance window, and an account of why this benchmark is hard to measure on at all, its folds being taxonomically non-exchangeable with a spread exceeding either addition on its own.

We therefore make placement the target of both the model and the measurement. The oracle we start from is DeepPeptide [Teufel et al., 2023b], the strongest general-purpose model for this formulation. It has three parts: a frozen protein language model, a CNN-BiLSTM encoder, and a linear-chain CRF [Lafferty et al., 2001]. The CRF has separate states for the start, the interior and the end of a segment, so any valid path must visit them in that order, and the decoder is boundary-aware by design. The encoder gives it nothing to work with: every state of the same segment type receives the same emission score. Closing that gap is the natural way to sharpen placement rather than merely find more segments, and we close it with two small, independent additions evaluated under the 5×4 nested cross-validation the benchmark itself uses, on the DeepPeptide dataset rebuilt from the 2026 UniProtKB/Swiss-Prot release (Sections 4.1 and 5).

2 Related work

Most predictors of proteolytic cleavage are specific by construction, covering a fixed set of proteases with known recognition motifs [Duckert et al., 2004, Song et al., 2012, Li et al., 2023, 2020a,b] or one product class in one kind of organism [Southey et al., 2006, Wang et al., 2024], and neither answers what an organism’s full complement of proteases does to a designed precursor. General-purpose models on frozen pLM embeddings [Lin et al., 2023, ESM Team, 2024] mostly score peptides already excised [Du et al., 2024, Zhu et al., 2025] rather than locating them, and PeptideLocator [Mooney et al., 2013] scores each residue rather than segmenting. DeepPeptide [Teufel et al., 2023b] is, to our knowledge, the only model that segments a precursor into typed peptide and propeptide spans, and it established the homology-partitioned benchmark used here [Teufel et al., 2023a], so we take both its architecture and its data pipeline as our starting point (Appendix A).

61 3 Method

62 **The base architecture.** DeepPeptide passes a frozen ESM-2 embedding through a CNN-BiLSTM
63 stack and decodes it with a linear-chain CRF [Lafferty et al., 2001] that expands each of
64 $\{\text{Peptide}, \text{Propeptide}\}$ into a chain of up to 50 position states, 101 in all, so a legal path can
65 only realize a contiguous segment of admissible length. The encoder computes one emission per
66 label and shares it across all 50 states of that label: the decoder is boundary-aware, the features
67 feeding it are not. The two additions close that gap from opposite ends of the pipeline (Figure 1b,
68 drawn at reading size in Appendix B). Other embedding sources and architectural modifications
69 were screened as well, with verdicts in Appendix F, and Appendix E sets out how the configuration
70 we train relates to the published one.

71 **Boundary head.** A small feed-forward block (LayerNorm, Linear, GELU, Linear, hidden size 64)
72 is applied to the contextual per-residue features. For each position and each segment type it emits
73 three numbers, scoring how much that position looks like a segment start, an interior position, or a
74 segment end, and these are added to the emissions of the corresponding position-fixed CRF states.
75 The final linear layer is initialized to zero, so training begins at exactly the base model and the head
76 only learns corrections that lower the loss. It costs 4,678 parameters on a trainable stack of 224,710.

77 **Adapter.** The second addition leaves the decoder alone and acts on the input. Before the per-
78 residue embedding reaches the CNN-BiLSTM it passes through LayerNorm, dropout, Linear,
79 GELU, dropout and a second LayerNorm, which reduces its width from 1280 to 256 for ESM-2.
80 A pLM embedding is trained on masked-residue recovery rather than cleavage-site localization, and
81 its feature distribution suits neither this task nor the small head that consumes it, so a trainable re-
82 projection on top of a frozen embedding is the standard remedy. It is the more expensive of the two,
83 a net +232,704 parameters.

84 4 Data and evaluation

85 4.1 Dataset

86 A precursor $x = (a_1, \dots, a_L)$ is labelled per residue with $y_t \in \{\text{None}, \text{Peptide}, \text{Propeptide}\}$
87 and contiguous runs of a label form typed segments, so the task is to recover which stretches of a
88 precursor become mature peptides and which become propeptides that are excised and discarded.
89 DeepPeptide built its dataset from PEPTIDE and PROPEP annotations in the 2022 Swiss-Prot release.
90 We rebuilt it with the same pipeline on the 2026 release [UniProt Consortium, 2025]. The collection
91 grows from 8,449 proteins to 9,619 (Table 2 in Appendix C), of which 8,897 both carry ESM-2
92 embeddings and enter the five folds used here. Folds come from GraphPart [Teufel et al., 2023a] at
93 a 30% pairwise-identity ceiling, balanced by cleavage-motif class. Segment-length filtering, motif
94 balancing and the full composition of the rebuild are given in Appendix C.

95 4.2 Evaluation criterion

96 Peptides and propeptides are matched separately. A prediction is a true positive when its start and its
97 end both fall within $\pm\tau$ residues of a true segment of the same type. Unmatched predictions are false
98 positives and unmatched true segments false negatives. Following the original evaluation we keep
99 $\tau = 3$ as the headline setting, since annotated sites carry a few residues of experimental uncertainty.
100 Sweeping τ down to zero then asks a second question, not whether a peptide was found but how
101 precisely its ends were placed. The reference implementation of this criterion carries an upstream
102 variable-shadowing bug that marks the wrong true segment as matched, so everything reported here
103 is re-scored with a corrected matcher (Appendix D).

104 **Agreement with the published baseline.** DeepPeptide reports precision 0.68 and recall 0.49 at
105 ± 3 , both segment types together, averaged over the twenty models of its nested cross-validation on
106 the 2022 data, which implies an F1 of 0.570. Running the base architecture on that release under the
107 same protocol scores 0.574 ± 0.069 with the original matcher and 0.590 ± 0.064 with the corrected
108 one of Appendix D. The F1 reproduces, the operating point does not: our original-matcher precision
109 is 0.614 and recall 0.541. The published system tunes its hyperparameters per outer fold, while

Table 1: 5×4 nested cross-validation with the corrected matcher. Each entry is the mean over the five outer folds, with the standard deviation across them as a subscript. Precision and recall are at the headline ± 3 tolerance, and the four F1 columns tighten it to an exact residue. *Growth* is how much a row’s F1 advantage over the ESM-2 base widens between ± 3 and an exact match.

Additions	at ± 3		F1 by tolerance				growth
	P	R	± 3	± 2	± 1	exact	
<i>ESM-2</i>							
none	0.621 $\pm$.025	0.539 $\pm$.045	0.576 $\pm$.029	0.544 $\pm$.027	0.467 $\pm$.020	0.354 $\pm$.010	—
boundary head	0.673 $\pm$.029	0.538 $\pm$.036	0.597 $\pm$.026	0.566 $\pm$.024	0.488 $\pm$.015	0.384 $\pm$.015	+0.010 $\pm$.011
adapter	0.628 $\pm$.019	0.579 $\pm$.040	0.602 $\pm$.026	0.572 $\pm$.024	0.491 $\pm$.011	0.396 $\pm$.020	+0.016 $\pm$.009
both	0.667 $\pm$.032	0.598 $\pm$.026	0.630 $\pm$.021	0.599 $\pm$.023	0.518 $\pm$.016	0.432 $\pm$.013	+0.023 $\pm$.019
<i>ESM-C 6B</i>							
none	0.620 $\pm$.017	0.560 $\pm$.027	0.588 $\pm$.016	0.558 $\pm$.013	0.483 $\pm$.007	0.371 $\pm$.013	+0.005 $\pm$.016
boundary head	0.731 $\pm$.010	0.572 $\pm$.030	0.641 $\pm$.022	0.608 $\pm$.026	0.532 $\pm$.035	0.429 $\pm$.030	+0.010 $\pm$.017
adapter	0.598 $\pm$.012	0.607 $\pm$.031	0.602 $\pm$.017	0.573 $\pm$.019	0.497 $\pm$.016	0.394 $\pm$.031	+0.014 $\pm$.023
both	0.690 $\pm$.013	0.646 $\pm$.033	0.666 $\pm$.018	0.636 $\pm$.015	0.554 $\pm$.016	0.466 $\pm$.021	+0.021 $\pm$.023

we hold one configuration fixed across all 20 cells so that every row of Table 1 is measured under the same setting (Appendix E). Under the corrected matcher throughout, the 2026 rebuild gives 0.576 ± 0.029 against that 0.590 ± 0.064 , so it reads 0.014 lower by a margin smaller than either run’s spread across folds.

5 Results

We ran the experiments in two stages.

Screening. We first split the data into seven GraphPart folds with fixed roles: four for training, one for epoch selection, and two held out, one for comparing architectures and one intended as a sealed test, which the screening numbers pool. Against that split we screened roughly a dozen modifications, with the protocol, the verdict table and the per-candidate figures in Appendix F, among them the segment-type trade-offs and data-scaling curves of Figures 7 and 8.

The folds of that split are not interchangeable (Figure 6 in that appendix), and under one assignment of roles several modifications changed the *sign* of their measured effect between the two held-out folds. A single draw of this kind cannot resolve an effect of 0.02–0.03, so none of those numbers is reported as a finding. It can, however, separate the consistently unhelpful from the worth paying for, and two cleared that bar: the boundary head and the adapter sat at or above the base throughout.

Confirmation. We then re-evaluated the two survivors with the 5×4 nested cross-validation DeepPeptide itself uses, on ESM-2 and on ESM-C 6B, so that the embedding is a factor of the design and not a third candidate. For each of five outer folds we train four models on three of the remaining folds and validate on the fourth, then test on the outer fold, which they never saw. An outer fold’s score is the mean over its four models, the estimate is the mean over the five outer folds, and the reported spread is the standard deviation across outer folds rather than across all 20 cells, since four models sharing an outer fold share a test set. Because selection happened at the screening stage, what follows is a test of two pre-specified hypotheses.

Table 1 reports the same 2×2 factorial on both embeddings under that protocol, with the corrected matcher.¹

The two additions act on different errors. On ESM-2 each improves F1 by a similar amount alone and together they give 0.054, against 0.079 on ESM-C, close to the sum of the parts in both cases. Splitting the ESM-2 rows into precision and recall shows the two are not competing for the same errors. The head is almost purely a precision effect, +0.052 against -0.001 of recall, the adapter mostly a recall effect, +0.040 against +0.007, and applied together they recover both, +0.046 and +0.059.

¹Implementation, the configuration behind every row, the per-cell scores every table is aggregated from and a script that rebuilds them: <https://anonymous.4open.science/r/cleavage-site-segmentation-5851>

142 **Tightening the tolerance.** Sweeping τ from ± 3 to an exact match, absolute F1 falls steeply for
143 every model, the strongest configuration from 0.666 to 0.466.

144 A model that is better overall retains a different fraction of its ± 3 score even when no boundary is
145 placed better, so we compare absolute gaps rather than ratios. Every gap an addition opens widens
146 as the tolerance tightens (the *growth* column of Table 1), and for each of those six the largest step
147 is the last, from ± 1 to exact. Only the embedding swap on its own does not grow. What separates
148 these models is whether they hit the exact residue.

149 A widening gap is not proof of better localization, since a model that finds more segments gains
150 ground at every tolerance. To hold detection fixed we compared each variant against the base on the
151 true segments that *both* localize within ± 3 , cell by cell. On that set the share of boundaries placed
152 on the exact residue is 0.630 for the base against 0.652 for the boundary head, 0.619 against 0.668
153 for the adapter and 0.621 against 0.698 for both, each positive on all five outer folds. A gate is itself
154 a choice, so Appendix G asks the same question with no window at all, and there the two blocks
155 part company. When a hit needs only a single residue of overlap the head covers less than the base,
156 recall -0.035 and below it on four folds of five, while gaining precision, and the adapter covers
157 more, $+0.024$ on four of five, while losing it. Their F1 there barely moves, -0.013 and $+0.006$,
158 because the components cancel. The head’s gain is placement alone, the adapter buys placement and
159 coverage together, and a ± 3 window records the mixture as recall. Scored at individual cleavage
160 sites the head’s gain lies entirely on the C-terminal side.

161 **Capacity the base decoder cannot use.** The base architecture on ESM-C 6B reaches $0.588 \pm$
162 0.016 against 0.576 ± 0.029 on ESM-2, a paired difference of $+0.012 \pm 0.019$ over the five outer
163 folds, so an embedding twice as wide buys no confirmed improvement on its own while either
164 addition on the narrower one does. What they are worth there depends on which one. Paired by outer
165 fold against its own base, the head gains $+0.054 \pm 0.026$ on ESM-C 6B against $+0.021 \pm 0.007$ on
166 ESM-2, the adapter $+0.014 \pm 0.014$ against $+0.026 \pm 0.018$, and the two together $+0.079 \pm 0.009$
167 against $+0.054 \pm 0.016$, all six positive on five folds of five, though the ESM-C adapter’s spread
168 reaches zero. The adapter re-projects an embedding trained for masked-residue recovery, and the
169 less mismatched it is the less there is to re-project, whereas the head supplies position-specific
170 evidence and a richer embedding carries more of it. Read the other way, the swap is worth $+0.012$
171 F1 under the base architecture and $+0.036$ under both additions, a split Appendix G reproduces on
172 the protein-cells common to all runs.

173 6 Conclusion and limitations

174 A zero-initialized boundary head at the decoder and a lightweight adapter at the input improve
175 segment F1 by $+0.054$ together on ESM-2 and $+0.079$ on ESM-C 6B. What matters for a design
176 loop is where that gain comes from, and the two blocks answer differently. At a gate asking only
177 for one residue of overlap the head covers less than the base while the adapter covers more, so the
178 head’s contribution is placement alone and the adapter’s placement on top of coverage. What grows
179 as the requirement tightens is placement, reaching $+0.077 \pm 0.005$ at exact match on ESM-2 and
180 $+0.095 \pm 0.015$ on ESM-C 6B, each against its own base. An oracle asked whether a construct
181 yields the peptide it was designed to yield is read at exactly that end of the curve, which is the end
182 the benchmark reports least well.

183 **The folds stay unequal.** Averaging over folds does not make them comparable, it only stops one
184 of them deciding the result. GraphPart keeps homologs together, and taxa *are* homology clusters:
185 313 of 714 *Conus* sequences fall in fold 1 against 16 in fold 2, all 293 *Cyriopagopus* avoid fold 0
186 entirely, and the folds differ in size by a factor of two (Table 3 in Appendix C). The base architecture
187 scores 0.572, 0.599, 0.530, 0.576 and 0.604 on outer folds 0 to 4, a range of 0.074, three and a half
188 times the $+0.021$ effect the same experiment resolves once averaged over them.

189 **What the protocol does not settle.** Nothing was selected on outer-fold scores inside confirmation,
190 since every cell of the 2×2 is reported, but the screening that chose the candidates ran on the same
191 proteins, and sealing that would take a third held-out level at roughly 1,600 GPU-hours for the grid,
192 on 2/5 of an already small dataset. The spread covers fold composition and not seed variation, and
193 the base row is untuned rather than the tuned published system (Appendix E).

References

- José Juan Almagro Armenteros, Marco Salvatore, Olof Emanuelsson, Ole Winther, Gunnar von Heijne, Arne Elofsson, and Henrik Nielsen. Detecting sequence signals in targeting peptides using deep learning. *Life Science Alliance*, 2(5):e201900429, 2019. doi: 10.26508/lsa.201900429.
- Zhenjiao Du, Xingjian Ding, William Hsu, Arslan Munir, Yixiang Xu, and Yonghui Li. pLM4ACE: a protein language model based predictor for antihypertensive peptide screening. *Food Chemistry*, 431:137162, 2024. doi: 10.1016/j.foodchem.2023.137162.
- Peter Duckert, Søren Brunak, and Nikolaj Blom. Prediction of proprotein convertase cleavage sites. *Protein Engineering, Design and Selection*, 17(1):107–112, 2004. doi: 10.1093/protein/gzh013.
- Ahmed Elnaggar, Hazem Essam, Wafaa Salah-Eldin, Walid Moustafa, Mohamed Elkerdawy, Charlotte Rochereau, and Burkhard Rost. Ankh: Optimized protein language model unlocks general-purpose modelling, 2023. URL <https://arxiv.org/abs/2301.06568>.
- ESM Team. Esm cambrian: Revealing the mysteries of proteins with unsupervised learning. <https://www.evolutionaryscale.ai/blog/esm-cambrian>, December 2024.
- Thomas Hayes, Roshan Rao, Halil Akin, Nicholas J. Sofroniew, Deniz Oktay, Zeming Lin, Robert Verkuil, Vincent Q. Tran, Jonathan Deaton, Marius Wiggert, Rohil Badkundri, Irhum Shafkat, Jun Gong, Alexander Derry, Raul S. Molina, Neil Thomas, Yousuf A. Khan, Chetan Mishra, Carolyn Kim, Liam J. Bartie, Matthew Nemeth, Patrick D. Hsu, Tom Sercu, Salvatore Candido, and Alexander Rives. Simulating 500 million years of evolution with a language model. *Science*, 387(6736):850–858, 2025. doi: 10.1126/science.ads0018.
- Michael Heinzinger, Konstantin Weissenow, Joaquin Gomez Sanchez, Adrian Henkel, Milot Mirdita, Martin Steinegger, and Burkhard Rost. Bilingual language model for protein sequence and structure. *NAR Genomics and Bioinformatics*, 6(4):lqae150, 2024. doi: 10.1093/nargab/lqae150.
- John D. Lafferty, Andrew McCallum, and Fernando C. N. Pereira. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Proceedings of the Eighteenth International Conference on Machine Learning (ICML)*, pages 282–289, 2001.
- Fuyi Li, Jinxiang Chen, André Leier, Tatiana Marquez-Lago, Quanzhong Liu, Yanze Wang, Jerico Revote, A. Ian Smith, Tatsuya Akutsu, Geoffrey I. Webb, Lukasz Kurgan, and Jiangning Song. DeepCleave: a deep learning predictor for caspase and matrix metalloprotease substrates and cleavage sites. *Bioinformatics*, 36(4):1057–1065, 2020a. doi: 10.1093/bioinformatics/btz721.
- Fuyi Li, André Leier, Quanzhong Liu, Yanan Wang, Dongxu Xiang, Tatsuya Akutsu, Geoffrey I. Webb, A. Ian Smith, Tatiana Marquez-Lago, Jian Li, and Jiangning Song. Procleave: predicting protease-specific substrate cleavage sites by combining sequence and structural information. *Genomics, Proteomics & Bioinformatics*, 18(1):52–64, 2020b. doi: 10.1016/j.gpb.2019.08.002.
- Fuyi Li, Cong Wang, Xudong Guo, Tatsuya Akutsu, Geoffrey I. Webb, Lachlan J. M. Coin, Lukasz Kurgan, and Jiangning Song. ProsperousPlus: a one-stop and comprehensive platform for accurate protease-specific substrate cleavage prediction and machine-learning model construction. *Briefings in Bioinformatics*, 24(6):bbad372, 2023. doi: 10.1093/bib/bbad372.
- Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023. doi: 10.1126/science.ade2574.
- Catherine Mooney, Niall J. Haslam, Gianluca Pollastri, and Denis C. Shields. Towards the improved discovery and design of functional peptides: common features of diverse classes permit generalized prediction of bioactivity. *PLoS ONE*, 7(10):e45012, 2012. doi: 10.1371/journal.pone.0045012.

242 Catherine Mooney, Niall J. Haslam, Thérèse A. Holton, Gianluca Pollastri, and Denis C. Shields.
 243 PeptideLocator: prediction of bioactive peptides in protein sequences. *Bioinformatics*, 29(9):
 244 1120–1126, 2013. doi: 10.1093/bioinformatics/btt103.

245 Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo,
 246 Myle Ott, C. Lawrence Zitnick, Jerry Ma, and Rob Fergus. Biological structure and function
 247 emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the*
 248 *National Academy of Sciences*, 118(15):e2016239118, 2021.

249 Jiangning Song, Hao Tan, Andrew J. Perry, Tatsuya Akutsu, Geoffrey I. Webb, James C. Whisstock,
 250 and Robert N. Pike. PROSPER: an integrated feature-based tool for predicting protease substrate
 251 cleavage sites. *PLoS ONE*, 7(11):e50300, 2012. doi: 10.1371/journal.pone.0050300.

252 Jiangning Song, Fuyi Li, André Leier, Tatiana T. Marquez-Lago, Tatsuya Akutsu, Gholamreza Haf-
 253 fari, Kuo-Chen Chou, Geoffrey I. Webb, and Robert N. Pike. PROSPERous: high-throughput
 254 prediction of substrate cleavage sites for 90 proteases with improved accuracy. *Bioinformatics*,
 255 34(4):684–687, 2018. doi: 10.1093/bioinformatics/btx670.

256 Bruce R. Southey, Andinet Amare, Tyler A. Zimmerman, Sandra L. Rodriguez-Zas, and Jonathan V.
 257 Sweedler. NeuroPred: a tool to predict cleavage sites in neuropeptide precursors and provide the
 258 masses of the resulting peptides. *Nucleic Acids Research*, 34(Web Server issue):W267–W272,
 259 2006. doi: 10.1093/nar/gkl161.

260 Felix Teufel, José Juan Almagro Armenteros, Alexander Rosenberg Johansen, Magnús Halldór Gís-
 261 lason, Silas Irby Pihl, Konstantinos D. Tsirigos, Ole Winther, Søren Brunak, Gunnar von Heijne,
 262 and Henrik Nielsen. SignalP 6.0 predicts all five types of signal peptides using protein language
 263 models. *Nature Biotechnology*, 40(7):1023–1025, 2022. doi: 10.1038/s41587-021-01156-3.

264 Felix Teufel, Magnús Halldór Gíslason, José Juan Almagro Armenteros, Alexander Rosenberg Jo-
 265 hansen, Ole Winther, and Henrik Nielsen. GraphPart: homology partitioning for biological se-
 266 quence analysis. *NAR Genomics and Bioinformatics*, 5(4):lqad088, 2023a. doi: 10.1093/nargab/
 267 lqad088.

268 Felix Teufel, Jan Christian Refsgaard, Christian Toft Madsen, Carsten Stahlhut, Mads Grønborg, Ole
 269 Winther, and Dennis Madsen. Deeppeptide predicts cleaved peptides in proteins using conditional
 270 random fields. *Bioinformatics*, 39(10):btad616, 10 2023b. ISSN 1367-4811. doi: 10.1093/
 271 bioinformatics/btad616. URL <https://doi.org/10.1093/bioinformatics/btad616>.

272 UniProt Consortium. UniProt: the universal protein knowledgebase in 2025. *Nucleic Acids Re-*
 273 *search*, 53(D1):D609–D617, 2025. doi: 10.1093/nar/gkae1010. URL <https://academic.oup.com/nar/article/53/D1/D609/7902999>.

275 Michel van Kempen, Stephanie S. Kim, Charlotte Tumescheit, Milot Mirdita, Jeongjae Lee,
 276 Cameron L. M. Gilchrist, Johannes Söding, and Martin Steinegger. Fast and accurate pro-
 277 tein structure search with Foldseek. *Nature Biotechnology*, 42:243–246, 2024. doi: 10.1038/
 278 s41587-023-01773-0.

279 Lei Wang, Chen Huang, Mingxia Wang, Zhidong Xue, and Yan Wang. NeuroPred-PLM: an inter-
 280 pretable and robust model for neuropeptide prediction by protein language model. *Briefings in*
 281 *Bioinformatics*, 24(2):bbad077, 2023. doi: 10.1093/bib/bbad077.

282 Lei Wang, Zilu Zeng, Zhidong Xue, and Yan Wang. DeepNeuropePred: a robust and universal tool
 283 to predict cleavage sites from neuropeptide precursors by protein language model. *Computational*
 284 *and Structural Biotechnology Journal*, 23:309–315, 2024. doi: 10.1016/j.csbj.2023.12.004.

285 Ming Zhang, Jianren Zhou, Xiaohua Wang, Xun Wang, and Fang Ge. DeepBP: ensemble deep
 286 learning strategy for bioactive peptide prediction. *BMC Bioinformatics*, 25(1):352, 2024. doi:
 287 10.1186/s12859-024-05974-5.

288 Lun Zhu, Hao Sun, and Sen Yang. BPFun: a deep learning framework for bioactive peptide function
 289 prediction using multi-label strategy by transformer-driven and sequence rich intrinsic informa-
 290 tion. *BMC Bioinformatics*, 26(1):187, 2025. doi: 10.1186/s12859-025-06190-5.

291 **Impact Statement**

292 This paper improves a computational tool for predicting proteolytic cleavage products from pro-
293 tein sequence. Potential applications include interpreting proteomics data, informing vaccine and
294 peptide-drug design, and studying disease-associated proteolysis. We are not aware of societal risks
295 specific to this work beyond those generally associated with more accurate protein sequence analysis
296 tools.

297 Appendix

298 A Extended review of prior work

299 Methods for predicting proteolytic peptides fall into three broad classes.

300 **Protease-specific models.** These predict cleavage only for a limited set of enzymes, most of which
301 have a known recognition motif. ProP [Duckert et al., 2004] targets cleavage by the PACE/PC family
302 in animals and plants. PROSPER [Song et al., 2012] and its successor PROSPERous [Song et al.,
303 2018] predict cleavage for one chosen enzyme at a time, as does the later ProsperousPlus [Li et al.,
304 2023]. DeepCleave [Li et al., 2020a] is restricted to caspases and matrix metalloproteinases, and
305 Procleave [Li et al., 2020b] likewise handles one enzyme at a time and additionally requires 3D
306 structural input. The shared limitation of this class is scope: in practice, one is usually interested in
307 the combined effect of all proteases active in an organism, not any single enzyme.

308 **Organism- or peptide-type-specific models.** A separate class targets neuropeptides or other bioac-
309 tive fragments of one product type. NeuroPred [Southey et al., 2006] and the species-agnostic
310 DeepNeuroPred [Wang et al., 2024] only detect cleavage near a handful of basic residues, while
311 NeuroPred-PLM [Wang et al., 2023] does not localize cleavage sites at all but classifies an already-
312 excised sequence. A related group of models finds signal subsequences rather than any bioactive
313 peptide: SignalP [Teufel et al., 2022] and TargetP [Almagro Armenteros et al., 2019] identify the
314 type of signal peptide, which indicates where a protein is trafficked.

315 **General-purpose models on pLM embeddings.** Protein language models (pLMs) are transform-
316 ers trained without supervision on large sequence corpora, typically by masked-residue recovery,
317 and they produce transferable per-residue embeddings: ESM-2 [Lin et al., 2023], the earlier ESM
318 line [Rives et al., 2021], ESM Cambrian [ESM Team, 2024, Hayes et al., 2025], and Ankh [El-
319 naggar et al., 2023]. Compact task-specific heads are then trained on top of these frozen embed-
320 dings. Besides **DeepPeptide** [Teufel et al., 2023b] itself, this recipe underlies DeepNeuroPred,
321 NeuroPred-PLM, and SignalP 6, as well as models that score already-excised peptides rather than
322 finding them: pLM4ACE [Du et al., 2024] predicts ACE-inhibitory activity, and BPFun [Zhu et al.,
323 2025] and DeepBP [Zhang et al., 2024] predict broader bioactivity, building on earlier pre-pLM
324 work such as PeptideRanker [Mooney et al., 2012]. Before pLMs, PeptideLocator [Mooney et al.,
325 2013] addressed a problem close to ours, but its output is a per-residue bioactivity score rather than
326 an explicit segmentation, and it was the main point of comparison in the original DeepPeptide paper.

327 **Baseline architecture: DeepPeptide.** Since we build directly on it throughout this paper (Sec-
328 tion 3), DeepPeptide’s architecture is worth describing in detail rather than treating as a black box.
329 It first encodes a precursor sequence with a frozen pLM (ESM-2), whose per-residue embedding is
330 never updated. This embedding is refined by a convolution, a single-layer bidirectional LSTM and
331 a second convolution. The published description does not fix the widths: they are searched with
332 Optuna inside the inner cross-validation loop. The configuration we train throughout uses 32 con-
333 volutional filters of width 3 and an LSTM hidden size of 64 per direction, giving 224,710 trainable
334 parameters for the base model. These features are consumed by a linear-chain CRF decoder whose
335 state set is extended well beyond the three raw labels {None, Peptide, Propeptide}: each of the
336 two positive labels is expanded into a chain of states long enough to represent segments up to a fixed
337 maximum length (101 states in total for the 50-residue cap used here, Section 4.1), so that a legal
338 path through the CRF can only realize a contiguous segment of valid length, with dedicated states
339 marking its first and last positions. This makes segmentation, rather than independent per-residue
340 classification, the object the model is directly optimized for – and it is also where the gap described
341 in Section 1 lives: every state in this extended set that shares a type receives the same emission from
342 the encoder, regardless of whether it marks the start, interior, or end of a segment.

343 **Baseline dataset for this task.** Beyond its architecture, DeepPeptide also established the data re-
344 source this line of work relies on: precursor sequences from UniProtKB/Swiss-Prot annotated with
345 PEPTIDE and PROPEP feature types, partitioned into homology-aware folds with GraphPart [Teufel
346 et al., 2023a] at a 30% pairwise-identity threshold and additionally balanced by cleavage-flanking
347 motif, then evaluated with the same family of nested cross-validation we adopt (Section 5). This
348 gave the field a standardized, already homology-controlled benchmark rather than an ad hoc collec-
349 tion of previously published peptide lists, and both the curated data and the construction pipeline
350 that produced it are public. Section 4.1 describes what we do with that resource: rebuilding it on a

current UniProtKB/Swiss-Prot release, which adds 1,170 proteins net and refreshes the annotations behind every segment.

Among these, DeepPeptide is, to our knowledge, the only model that poses the task as full segmentation of the precursor into typed peptide and propeptide segments, and it is the strongest available baseline for this formulation. We therefore use its architecture and its dataset-construction methodology as the starting point for this work, rather than treating it as a system to audit: our contribution is an architectural addition to this general recipe (Section 3), evaluated under the same 5×4 nested cross-validation it uses (Section 4) on data rebuilt from a current release (Section 4.1).

B The architecture at reading size




Figure 2: The base architecture and the two additions, at reading size. Figure 1b shows the same diagram at the scale of a single panel, where it locates the two blocks rather than documenting them. Top: the pipeline, with the adapter between the frozen pLM and the CNN-BiLSTM encoder and the boundary head reading the encoder’s output. Bottom: what the head adds. The base model computes one emission per label and shares it across all 50 position states of that label, so the CRF’s start, interior and end states are fed identical evidence. The head emits three numbers per position and per segment type and adds them to the states they describe.

C Dataset details

Table 2: Dataset composition: the 2022 release used by the original DeepPeptide against the 2026 release rebuilt here. Counts are before homology partitioning, which leaves 7,623 proteins of the 2022 release and 8,994 of the 2026 one, of which 8,897 also carry an ESM-2 embedding and form the five folds used here.

	2022	2026
Proteins	8449	9619
Peptide segments	6372	7431
Propeptide segments	8211	9140
Median peptide length (residues)	21	20

Two filters from the original pipeline are kept unchanged. Segments shorter than 5 or longer than 50 residues are dropped: the CRF’s state count fixes the upper bound, and a segment shorter than five residues is barely scoreable at a ± 3 tolerance. Folds are then balanced by cleavage-motif class, obtained by k -means ($k = 50$) on ESM-2 embeddings of the four residues flanking each annotated boundary, on top of the 30% pairwise-identity ceiling GraphPart enforces between folds.




Figure 3: Composition of the rebuilt 2026 dataset: segment length, type, and genus distributions.

Table 3: Proteins per outer fold, and how three frequent genera distribute across them. *Cyriopagopus* and *Lycosa* are spiders, *Conus* a cone snail, and each is a homology cluster that GraphPart is obliged to keep intact. *Homo*, by contrast, is spread evenly (71/81/66/130/84).

	total	f0	f1	f2	f3	f4
All proteins	8897	1558	2572	1263	2025	1479
<i>Conus</i>	714	160	313	16	149	76
<i>Cyriopagopus</i>	293	0	62	45	9	177
<i>Lycosa</i>	163	5	42	10	106	0




Figure 4: Segment F1 at ± 3 on each outer fold, every configuration on one axis, with the band between the best and worst configuration shaded. The lines move together: fold 2 is the hardest for all eight and folds 1 and 4 the easiest, so most of the vertical movement within a line is the fold rather than the architecture. Within one configuration the score spans 0.040 to 0.074 across the folds, more than either addition is worth on its own, which is what makes a single split unable to resolve them and pairing by fold necessary. The ordering of configurations is largely preserved from fold to fold, which is why the paired comparisons of Section 5 are stable where the absolute levels are not.




Figure 5: Why a fold’s length profile is not a cosmetic difference. Segment F1 at ± 3 against the length of the true segment, with the number of true segments in each bin below. Quality is not flat in length: the base architecture scores 0.715 on segments of 25–29 residues and 0.274 on those of 45–50, a spread of 0.441, which is several times any effect this paper measures. A fold holding more long segments is therefore harder before any model touches it, and the folds do differ in length composition (Figure 6). This is what makes them non-exchangeable in the sense Section 6 means. The two additions help most where the base is already strongest and least at the long end, so they do not flatten the curve.

Table 4: The same scores as numbers: segment F1 at ± 3 per outer fold, mean $\pm$ standard deviation across the five.

Configuration	f0	f1	f2	f3	f4	mean
ESM-2, base	0.572	0.599	0.530	0.576	0.604	0.576 $\pm$ 0.029
ESM-2 + boundary head	0.593	0.613	0.560	0.589	0.629	0.597 $\pm$ 0.026
ESM-2 + adapter	0.622	0.626	0.561	0.595	0.606	0.602 $\pm$ 0.026
ESM-2 + both	0.639	0.650	0.603	0.613	0.646	0.630 $\pm$ 0.021
ESM-C 6B, base	0.575	0.611	0.571	0.590	0.592	0.588 $\pm$ 0.016
ESM-C 6B + boundary head	0.664	0.634	0.609	0.639	0.661	0.641 $\pm$ 0.022
ESM-C 6B + adapter	0.613	0.618	0.576	0.604	0.597	0.602 $\pm$ 0.017
ESM-C 6B + both	0.664	0.690	0.643	0.658	0.676	0.666 $\pm$ 0.018

366 D Metric implementation bug

367 While auditing the DeepPeptide implementation of the matching criterion (Section 4.2), we found
 368 a bug in the function that implements it (`get_counts_for_protein`), reproduced below in simplified form:
 369

```

370 for idx, row in true_df.iterrows(): # idx: true-segment index
371     true_start, true_stop = row['start'], row['stop']
372     for idx, row in pred_df.iterrows(): # idx overwritten with pred index
373         pred_start, pred_stop = row['start'], row['stop']
374         ...
375     if start_match and stop_match:
376         true_df.loc[idx, 'matched'] = True # BUG: idx is now the
377         pred_df.loc[idx, 'matched'] = True # predicted index, not the true one
378     break
379
```

381 The outer loop iterates over true segments and the inner loop over predicted segments, but both
 382 reuse the loop variable `idx`. By the time a match is found, `idx` refers to the predicted segment,
 383 so `true_df.loc[idx, 'matched']` marks the wrong true segment. The effect is not large on
 384 average (it is partly self-correcting), but it costs a true positive whenever it fires, so both recall and
 385 precision come out low, and F1 with them. Across the nested-CV grid the correction raises recall
 386 by 0.021 to 0.032 and precision by 0.005 to 0.010. The loss concentrates in proteins where the
 387 model predicts more segments than the protein actually has. This bug is present in the upstream
 388 DeepPeptide code, not something we introduced.

389 We did not silently patch this function in place, since that would shift absolute numbers without
 390 documenting it. Instead we implemented a separate, corrected matcher and used it to re-score the
 391 results in this paper: the earlier single-split runs reported in Appendix F, except two whose model
 392 classes no longer exist in the code and whose checkpoints cannot be rebuilt, and the full 5×4
 393 nested-cross-validation grid in Table 1 (all 160 cells of the eight configurations, and the 20 further
 394 cells of the 2022-release reproduction), by re-running test-partition inference from each cell’s saved
 395 checkpoint on the machine that held them. Every table in the main text uses the corrected matcher,
 396 and both readings of every cell are released alongside the code, with a script that rebuilds the paper’s
 397 tables from them.

398 Re-scoring the grid also recomputed the original matcher on the same decoded segments, so each
 399 cell could be checked against the number already recorded for it. Over the 160 cells for which both
 400 readings exist, the two agree to 1.2×10^{-3} F1 at worst and to 3×10^{-7} at the median, so what
 401 remains is inference nondeterminism rather than disagreement between two deterministic functions
 402 of the same segments.

403 Across configurations the correction shifts mean F1 by +0.015 to +0.022, always upward, since
 404 the bug was discarding true positives. On ESM-2 it leaves the ordering of the four configurations
 405 intact, base < boundary head < adapter < combined, before and after correction. That ordering
 406 is not robust and we build no argument on it: the boundary-head/adaptor gap is 0.005 against a
 407 fold-level standard deviation of 0.026, and per outer fold the full ordering holds in four folds of five
 408 after correction and three of five before it. On ESM-C 6B the ordering differs, base < adapter <
 409 boundary head < combined, the two middle configurations swapping places.

410 E Relation to the published DeepPeptide system

411 Our base architecture is DeepPeptide’s architecture, but it is not DeepPeptide’s published model,
 412 and the difference is worth stating because it is the reason Section 4.2 reproduces the published F1
 413 without reproducing the published precision and recall.

414 The released implementation spends the inner loop of its nested cross-validation on a hyperparam-
 415 eter search, and the twenty checkpoints behind the published numbers carry a different winning
 416 configuration for each outer fold: learning rate from 3.3×10^{-4} to 5.5×10^{-3} , dropout from 0.03
 417 to 0.69, convolutional dropout from 0.04 to 0.51, kernel size 3 or 5, 48 to 96 filters, hidden size 32
 418 or 48, and batch size 20 to 90. We hold a single configuration fixed across all twenty cells at the
 419 implementation’s defaults (learning rate 10^{-4} , dropout and convolutional dropout 0.1, kernel size 3,
 420 32 filters, hidden size 64, seed 42), with batch size 48 rather than the default 100, and spend the inner
 421 loop on epoch selection alone. The two added blocks carry settings of their own, chosen once at the

422 screening stage and then held fixed as well: the adapter projects to 256 channels with a dropout of
423 0.4, and the boundary head uses a hidden size of 64.

424 Our settings are the implementation’s defaults and they sit outside the range the published search
425 explored: a learning rate below all five of its ESM-2 winners, fewer convolutional filters than any
426 of them, and a wider recurrent hidden state than any of them. The base row is therefore a different
427 model from the published one, not a worse-tuned copy of it, which is what the precision-recall
428 difference reflects.

429 That choice is deliberate. A factorial is only interpretable if its cells differ in the factor under study
430 and in nothing else, and per-fold tuning would confound each addition with whatever the search
431 happened to find for it. Two consequences follow, and they are not symmetric between the two
432 blocks.

433 The boundary head is protected by its initialization. Its final linear layer starts at zero, so at the
434 first step the model with the head *is* the base model, and the head’s 4,678 parameters can only
435 add corrections that lower the loss from there. Whatever setting suits the base is by construction
436 a legitimate setting for the head, which cannot be handicapped relative to the base it is added to.
437 The adapter has no such protection. It changes the input width from 1280 to 256 and adds 232,704
438 parameters, roughly doubling the trainable stack, and a model of that shape may well want a different
439 learning rate and dropout from the one we hold fixed. Its measured effect, +0.026 on ESM-2 and
440 +0.014 on ESM-C 6B, should therefore be read as a lower bound.

441 The direction of this bias is the useful part: holding hyperparameters fixed can hide an effect but
442 cannot manufacture one, so every gain in Table 1 is conservative. What it does not license is the
443 reading that an untuned base makes the additions look better than they are. On the 2022 release our
444 untuned base reaches an F1 of 0.574 against the tuned published model’s implied 0.570 (Section 4.2),
445 so the base we improve on is not a weak one.

446 Two further notes on the reference implementation. Its recall guard tests the wrong denominator,
447 returning zero when no segment is predicted but dividing by the true-segment count that the guard
448 never checked. On this data the branch is unreachable, so the fault is latent rather than active. And
449 its evaluation script builds the ground truth with overlapping annotations kept separate while its
450 training dataset merges them, a discrepancy we did not attempt to reconcile because our scoring
451 rebuilds the ground truth from the labelled sequences directly.

452 **F Screening protocol and verdicts**

453 **F.1 Protocol**

454 Data was split into seven GraphPart folds (30% identity threshold, motif-balanced as in Section 4.1),
455 with folds assigned fixed roles: four folds for training, one for validation (epoch selection), one
456 for model selection (comparing architectures), and one originally intended as a fully sealed test
457 fold. This separated architecture selection from evaluation, but the evaluation itself was still a single
458 draw: one random assignment of roles, one pair of held-out folds. Several modifications changed
459 the sign of their measured effect between the two held-out folds. Replacing ESM-2 with ESM-C 6B
460 measured +0.074 on one and −0.011 on the other, ESM-C 6B against ESM-C 600M gave +0.028
461 and −0.021, and the net effect of the 3Di channel gave +0.015 and −0.018. Effects that kept their
462 sign still moved by a lot: the isolated sequence adapter measured −0.001 on one fold and +0.045
463 on the other. Figure 6 shows one reason, namely that the folds differ substantially in segment-length
464 composition. The two held-out folds turn out to be complementary in segment-length composition,
465 fold 2 concentrated at 10–24 residues and fold 5 bimodal, so results below pool them (model-select
466 $\cup$ sealed, $\approx$ 2,300 proteins) to reduce (but not eliminate) this instability. This pooling is why the
467 “sealed” fold is not, in practice, a held-out test set independent of model selection.

468 **F.2 Summary of tested modifications**

469 Table 5 summarizes the modifications tested under screening protocol, with the corrected matcher
470 (Appendix D) applied throughout. Two of these rows were carried forward as candidates, the bound-
471 ary head and the adapter (Section 3), and both were confirmed. The embedding swap was not a
472 candidate but became a factor of the confirmation design, and it is the row this table reads least well:




Figure 6: Why the folds of the earlier seven-fold split are not interchangeable: (a) the segment-length profile of each fold against the profile of the whole dataset, (b) the L_1 distance between the two. Fold 4 tracks the overall distribution closely (0.07) while five of the seven folds sit above 0.23. Balancing was done on cleavage motif and homology, not on segment length, so this axis was left free. The same holds for the five-fold split used in the main text, where the imbalance is taxonomic (Table 3).

473 nested cross-validation puts ESM-C 6B at 0.588 ± 0.016 against ESM-2’s 0.576 ± 0.029 (Table 1),
 474 overlapping intervals, where this table records a confident $+0.03$. Everything else below remains at
 475 the confidence of a single pooled split.

476 Note the apparent tension with Table 1: under this earlier protocol, the boundary head on ESM-2
 477 looked like it had no effect, while nested cross-validation resolved a confirmed $+0.021$ F1 gain
 478 for the same modification on the same embedding. We read this as evidence that the single-split
 479 protocol lacked the statistical power to resolve an effect of this size, not as a contradiction – and as
 480 the clearest illustration of why we do not treat single-split numbers as findings in this paper. The
 481 interaction these rows suggest, that the boundary head is worth more on the richer embedding, is
 482 not established by them: it is established under nested cross-validation in Section 5, where the same
 483 comparison is paired by outer fold and positive on five folds of five. What this table contributes is
 484 the earlier and much noisier version of it.

485 F.3 Additional exploratory modifications

486 A larger set of ideas was tried under the same single-split protocol and did not show a clear, con-
 487 sistent benefit. We list them for completeness and as candidates for future re-testing, without confi-
 488 dence intervals:




Figure 7: Precision/recall trade-offs by segment type when adding the 3Di channel or the bond-prediction loss, single-split protocol, pooled held-out folds.

Table 5: Modifications tested under the screening single-split protocol, with their measured effect on F1 (pooled held-out folds) and an informal verdict. None of these effect sizes should be compared directly to the nested-CV numbers in Table 1.

Modification	Type	Measured effect on F1	Verdict
ESM-C 6B instead of ESM-2	embedding swap	+0.03, unstable across folds	helps
Boundary head on ESM-C 6B	decoder addition	+0.05, +0.07, consistent	helps
Boundary head on ESM-2	decoder addition	≈ 0 (CI includes zero)	no effect
Structural channel (3Di)	extra input	propep. +0.02 / pep. -0.03, net trade-off	no effect
ESM-C 6B compression 2560→256	optimization	≈ 0 , 10× narrower input and 16× fewer trainable parameters	no effect
Sequence adapter on ESM-C 6B (pLM re-projection)	input adapter	+0.022, CI [+0.008, +0.038]	helps
Bond-prediction auxiliary loss (on ESM-C 6B)	extra loss	pep. -0.05, net -0.015, neutral on ESM-2	harmful
Telescopic segment CRF	decoder addition	≈ 0 on ESM-2, +0.022 on ESM-C 6B, no CI available	unresolved

- Known-peptide dictionary.** An Aho–Corasick index of peptides from the training data was used to flag substring matches in a query sequence, injected as a bonus on the CRF state emissions, as a bias on the start, inside and end emissions specifically, fused into the encoder hidden state, and concatenated with the embedding before the encoder. No consistent benefit was observed, and the checkpoint for one of the variants can no longer be loaded.
- Structural features.** Features extracted from predicted 3D structures (AlphaFold2 representations via the AFToolkit codebase, with several confidence-based filtering variants) and from the ProstT5 [Heinzinger et al., 2024] structural alphabet (3Di, as used in Foldseek [van Kempen et al., 2024]) were tested as additional input channels, and results were mixed (Table 5).
- LoRA fine-tuning.** Low-rank adaptation of the last few pLM layers, instead of fully frozen embeddings, scored 0.558 and 0.567 against 0.621 for the frozen baseline, all three on the 2022 release under its own five-fold split rather than the protocol of Appendix F. It also ran for fewer epochs in the same time budget, though we did not measure the cost directly.

- **Projector variants.** Multi-scale and multi-branch projectors between the embedding and the CNN-BiLSTM were tried as alternative re-projection schemes to the adapter of Section 3.
- **Structural-projection width and telescopic CRF.** The width of the structural-feature projection (16/32/48 units) left pooled F1 flat at 0.691, 0.697 and 0.693, the telescopic segment CRF, an alternative decoder formulation that scores whole segments through a relative-position head, came out level with the base decoder on ESM-2 and +0.022 ahead on ESM-C 6B. Per-protein outputs were not retained for either run, so neither number carries a confidence interval and we treat the modification as untested rather than as having no effect.

None of these modifications showed a consistent gain large enough, under the single-split protocol, to justify testing under nested cross-validation ahead of the two reported in the main text. The complete experiment log (including runs not summarized above) is maintained in the project repository rather than reproduced here, since it was collected under a protocol we no longer treat as sufficient evidence on its own.





Figure 8: F1 (± 3) against the number of training proteins, single-split protocol. ESM-2 rises from 0.498 at 40% of the training folds to 0.583 at 85%, then falls back to 0.572 at 100%, so it is still gaining over most of the range but not at the last point. The ESM-C 6B curves are flat from the smallest size tested, with movement smaller than the bootstrap interval, so nothing is observed saturating: they simply never climb.


F.4 What the screening runs looked like

Three figures from the screening stage are worth keeping, each for a different reason, and none of them should be read as confirmed evidence in the sense of Section 5. Figure 9 is the comparison the verdict column of Table 5 summarizes, with the bootstrap intervals that made most of those verdicts unsafe. Figure 10 is the earliest form of the interaction that Section 5 later confirms under nested cross-validation, and shows that the boundary head was already worth more on the wider embedding a protocol ago. Figure 11 asks what recall actually tracks: for both ESM-C models it follows the maximum sequence identity between a held-out peptide and any training segment far more closely than it follows how many training segments the protein’s genus contributes, which is the behaviour of a model retrieving near neighbours rather than one that has learned the cleavage grammar. On the ESM-2 baseline the two axes are closer to comparable. This is a single split and we draw no conclusion from it, but it is the observation that would most repay a proper test.



hollow markers — gated-model controls, not rungs of the ladder
dashed rule — gated adapter: +0.022 (isolated by ablation),
significant on average, but the gain is mostly on fold 5

Figure 9: Screening comparison on the pooled held-out folds at ± 3 : F1, precision and recall with bootstrap confidence intervals over 2,322 proteins. The top row is a gated-projector control at full width, not the configuration carried into the main text.



error bars — 95% percentile bootstrap CI, resampling proteins (not folds)

Figure 10: Boundary head $\times$ embedding interaction under the screening protocol: F1 gain from adding the boundary head on top of ESM-2 against on top of ESM-C 6B. This is the earlier, unconfirmed version of the interaction that Section 5 measures under nested cross-validation.




Figure 11: Recall on held-out peptides as a function of (a) how well-represented the protein's genus is in training and (b) maximum sequence identity to a training segment, screening protocol. For the two ESM-C models recall tracks maximum identity to a training segment far more than genus abundance, while for the ESM-2 baseline the two axes are closer to comparable. The x axis in (a) counts training *segments* in the genus, not proteins.

G Boundary placement beyond the ± 3 window

Section 5 makes two claims that no single tolerance can settle: that the additions place boundaries better rather than merely finding more segments, and that what they buy survives a criterion with no acceptance window at all. This appendix reports the evidence for both.

Holding detection fixed. The comparison in Section 5 is paired. For each cell we take the true segments that the variant and the base both localize within ± 3 , so both models have found the segment and only its placement is at stake, and ask how often each puts a boundary on the annotated residue. A configuration is twenty cells, five outer folds by four inner ones, and the four cells sharing an outer fold are all tested on it, so each of the 15,348 true segments in the split is scored four times: 61,392 segment-instances in all, of which the base localizes 32,326. Pooled over the twenty cells the paired set holds 27,013 instances for the boundary head, 29,277 for the adapter and 29,394 for the two together. These are instances rather than distinct segments, so the paired comparison rests on about a quarter as many independent segments as the counts suggest.

What each block does to the count. Those denominators are themselves the result. Of the 32,326 instances the base localizes, the boundary head localizes 320 fewer and the adapter 2,236 more, with both together 3,289 more. About 1,400 of the adapter’s gain is coverage, its recall at a one-residue gate being $+0.024$ over the base on these 61,392 instances, and the rest is segments the base already overlapped but placed outside the window. The head therefore improves placement while covering marginally less, which is the precision effect of Table 1 seen segment by segment, and the adapter covers more while placing better still. On the paired set the head moves 13.2% of boundaries closer to the annotated residue and 11.4% further away, a near-even trade that nets out positive, while the two together move 15.0% closer against 7.5% further.

Dropping the window. A gate is itself a choice, so we also score the same predictions with no tolerance. Predicted and true segments of the same type are matched greedily one to one by intersection over union, and a match counts as a hit when the overlap clears a threshold on intersection over union. Precision and recall use all predicted and all true segments as their denominators, so nothing is conditioned on what either model happened to find. Every configuration is scored on the 35,504 protein-cells common to all runs, which takes the coverage difference between the two embeddings (8,897 against 8,999 proteins) out of the comparison. On that common set the ESM-2 to ESM-C 6B swap is worth $+0.005$ F1 at $\text{IoU} \geq 0.5$ under the base architecture, positive on two outer folds of five, and $+0.024$ under both additions, on five of five, reproducing the split that Section 5 measures at ± 3 .

Table 6: Segment F1 under overlap matching. The first column requires only that a predicted and a true segment of the same type share one residue, and the rest require an intersection over union of at least τ . Same cells as Table 1, restricted to the protein set common to every run, mean over the five outer folds.

	overlap	$\text{IoU} \geq 0.5$	$\text{IoU} \geq 0.8$	$\text{IoU} \geq 0.9$	$\text{IoU} \geq 0.95$
<i>ESM-2</i>					
base	0.729	0.667	0.569	0.487	0.411
+ boundary head	0.715	0.668	0.585	0.515	0.438
+ adapter	0.734	0.686	0.590	0.516	0.441
+ both	0.746	0.703	0.616	0.546	0.474
<i>ESM-C 6B</i>					
base	0.738	0.672	0.574	0.496	0.419
+ boundary head	0.735	0.704	0.624	0.555	0.478
+ adapter	0.729	0.681	0.587	0.514	0.440
+ both	0.759	0.727	0.649	0.582	0.507

Mostly placement. The looser the criterion, the less either addition is worth. Paired by outer fold against the ESM-2 base, at a gate that asks only for a single residue of overlap the boundary head is worth -0.013 ± 0.026 and the adapter $+0.006 \pm 0.012$, neither distinguishable from zero. At $\text{IoU} \geq 0.5$ they are worth $+0.001 \pm 0.021$ and $+0.019 \pm 0.017$, and at $\text{IoU} \geq 0.9$ $+0.028 \pm 0.004$ and $+0.028 \pm 0.015$, both on five folds of five. Underneath that near-zero F1 at the overlap gate

568 the two move in opposite directions: the head issues 219 fewer segments per fold, buying +0.022
569 precision for -0.035 recall, while the adapter issues 130 more and buys +0.024 recall for -0.021
570 precision. The head finds no more of the annotation than the base does and the adapter only modestly
571 more. What changes most is where the boundaries land, which is the conclusion Table 1 reaches
572 through a tolerance, reached here without one.

573 The mean IoU taken over *all* true segments moves the other way for the head, 0.576 to 0.559, while
574 its F1 rises at every threshold from 0.5 up. That is the precision effect of Table 1 seen from another
575 angle: the head withdraws weak predictions, so some true segments lose a poor match altogether,
576 and the ones that keep a match are placed better.

577 **Which end moves.** Scoring cleavage sites rather than segments, at exact placement and against
578 all annotated sites, separates the two additions again. The boundary head improves the C-side of a
579 segment, $+0.028 \pm 0.006$ on five folds of five, and leaves the N-side alone, -0.007 ± 0.020 and
580 positive on only two of five. The adapter improves both, the N-side by $+0.022 \pm 0.006$ and the C-
581 side by $+0.028 \pm 0.009$, and the two together by $+0.042 \pm 0.010$ and $+0.060 \pm 0.006$ respectively.
582 Split by segment type, the head's C-side gain is largest on propeptides ($+0.037 \pm 0.021$, five of five),
583 whose C-terminal cut is where the mature peptide begins and is the weaker of the two C-sides: the
584 base model scores 0.476 there against 0.525 on peptide C-sides.